

Bitmap Index

Objectives:

- Basic introduction to a very different indexing concept.
- Place for the method in application use and culture

Reading: text 14.7

Thanks: a few slides borrowed from Ramakrishnan text slides

10BitVectorIndex10

Database Engineering

1

Review & Backfill from last time

10BitVectorIndex10

Database Engineering

2

Access-Path

Access-Path:

- A method for fetching data from disk.
- or
- The sequence of blocks/pages required to located and retrieve data.

Usage: "Choose an access path",
e.g. either a sequential scan based on the primary index, or an available secondary index.

12: plan trees

Database Management & Engineering

3

RDBMS Storage


Default: Two Access Paths

B+ tree: Primary index

- Search key aka index key
= primary key

Table Scan aka sequential access path

- sorted on primary key
- (if possible) contiguously on disk.
 - Fully contiguous = 100% clustered



15SmartIndexing

Biological Data Models

4

Cultural Background (caveat)

Bitmap index:

- outcome of
 - research in parallel database systems
 - development and research in decision support systems (OLAP, data warehouses)
- Viewed a *very* inconsistent with set theoretic models of database - (even just a hack)
 - academics little textbook material
 - actual performance benefit in many circumstances leaves no choice
 - by serendipity, works well in circumstances much broader than first promulgated.

10BitVectorIndex10

Database Engineering

5

Why?

Speculation [Miranker]

- In a manner consistent enough with relational databases.
- Create a tiny, high density encoding of a database.
 - tiny often means fast.
 - I/O: definitely
 - CPU:

10BitVectorIndex10

Database Engineering

6

Effectiveness → Taking over

Big Data in the relational model (used directly)

→ Column Store Databases // next time

10BitVectorIndex10

Database Engineering

7

Bitmaps, as introduction to *Bloom Filters*

- “Bloom Filter” primary method of indexing in Big Data (Hadoop, but others as well)

10BitVectorIndex10

Database Engineering

8

Basic Idea (1)

- For a select (few) columns/attributes

Male: row1, row3, row5, row6...

Female : row2, row4, ...

- represent the contents of each row of the inverted index as a bitmap.

Male: 101011...

Female:010100...

10BitVectorIndex10

Database Engineering

9

Expandable for Any Type

Faculty Table

RowId	FacSSN	...	FacRank
1	098-55-1234		Asst
2	123-45-6789		Asst
3	456-89-1243		Asse
4	111-09-0245		Prof
5	931-99-2034		Asst
6	998-00-1245		Prof
7	287-44-3341		Asse
8	230-21-9432		Asst
9	321-44-5588		Prof
10	443-22-3356		Asse
11	559-87-3211		Prof
12	220-44-5688		Asst

Bitmap Index on FacRank

	Bitmap
Asst	110010010001
Asse	001000100100
Prof	000101001010

10BitVectorIndex10

Database Engineering

10

Notes:

- As the number of values increases the bitmaps get bigger
 - > results that suggest this method is of value for small domains (NOT, more later)
 - Even continuous data types (numbers and other infinitely sized) this works

10BitVectorIndex10

Database Engineering

11

Defining Bitmap Indexes

- Oracle syntax:
 - CREATE BITMAP INDEX
ON Faculty(FacRank);
- PostgreSQL, MS SQL Server, IBM DB2 – users can't declare
 - So effective, query engine will
 - build an index at the beginning of a query,
 - Throw it out when query is done

10BitVectorIndex10

Database Engineering

12

Bit vectors can also be used to represent join indexes

- Join indexes:
 - like a "join table" for a many-to-many association
- Oracle syntax:
 - CREATE BITMAP INDEX
 - ON Sales (Customer.State)
 - FROM Sales, Customer
 - WHERE Sales.Customer_key = Customer.Customer_key
- Creates an index with <State, Sales RID> entries
- (revisit this at the end of the semester as part of parallel DBMS)

10BitVectorIndex10

Database Engineering

13

How are these useful?

- Select *
- From Faculty
- Where Faculty.gender = "F" and
Faculty.FacRank = "asst"
- Output records determined by ANDing the two bitmaps together.
 - indexes for all columns of a table are the same length

10BitVectorIndex10

Database Engineering

14

Example

Select *

From Faculty

Where Faculty.gender = "F" and
Faculty.FacRank = "asst"

	Bitmap
Asst	110010010001
F	01010...
result	01000...

10BitVectorIndex10

Database Engineering

15

Sizing:

- Consider 8 million rows of faculty:
 - five megabytes of index
 - two megabytes of I/O to process the query based on the index.
 - _____ megabytes of I/O without the index.
 - Will another index work?

10BitVectorIndex10

Database Engineering

16

- What is the impact on size of the index?
 - as a function of number of
 - distinct values, m
 - rows, n

--> (initially thought:) good for small enumerated types

10BitVectorIndex10

Database Engineering

17

Idea (2a) Compression

- Suppose we have a large domain,
 - Example, domain size = 20
 - probability that a bit is 1?
 - how many zero's in a row?
- (draw example on board)

10BitVectorIndex10

Database Engineering

18

Idea (2a) Compression

- The bitmaps are compressed
- Run length encoding methods,
 - (run length of zero's = number of zero's, terminated by a one)
 - e.g. 001,000,001,0000
 - Run length of, 2, followed by run length of 5

What about zeros at the end of the vector?

Answer: Know the total number of rows....

10BitVectorIndex10

Database Engineering

19

How to represent run lengths?

simple sequence of binary numbers won't work

Consider,

e.g. 001,000,001,000,0

Run length of, 2, followed by run length of 5

10, 101 = 10101

But 10101, could be run length, 1010, 1

i.e. run length 10 followed by 1

→ 000,000,000,010,1

10BitVectorIndex10

Database Engineering

20

One Method: Exploits Unary Encoding

Unary coding

From Wikipedia, the free encyclopedia

Unary coding, sometimes called **thermometer code**, is a number is understood as *non-negative integer*) or with *n* - is represented as 111110 or 11110. Some representation generally. Unary coding is both a *Prefix-free code* and a *!*

n (non-negative)	n (strictly positive)	Unary code
0	1	0
1	2	10
2	3	110
3	4	1110

10BitVectorIndex10

Database Engineering

21

One Method: Exploits Unary Encoding

- Given, e.g. 001,000,001,000,0

Run length of, 2, followed by run length of 5

For each run, represent pair (x : y)

x, number of bits in unary to represent length of the run in binary
y, length of the run in binary

For run length 2

x = 10

y = 10, x:y 1010

10BitVectorIndex10

Database Engineering

22

One Method: Unary Encoding

- Given, e.g. 001,000,001,000,0

Run length of, 2, followed by run length of 5

For each run, represent pair (x : y)

x, number of bits in unary to represent length of the run
y, length of the run in binary

For run length 2

x = 10

y = 10, x:y 1010

For run length 5

x = 110

y = 101 x:y 110101

10BitVectorIndex10

Database Engineering

23

One Method: Unary Encoding

- Given, e.g. 001,000,001,000,0

Run length of, 2, followed by run length of 5

For run length 2

x = 10

y = 10, x:y 1010

For run length 5

x = 110

y = 101 x:y 110101

001,000,001,000,0 → 1010 110101

10BitVectorIndex10

Database Engineering

24

Decoding

Decode 1101001011

locate length of first run by finding "0"

1101001011 → run length takes 3 bits to represent

110, 100 (x/3, y/4) → 00001

1011 remains → run length takes 2 bits to represent

10, 11 (x/2, y/3) → 0001

so

1101001011 → 00001 0001

10BitVectorIndex10

Database Engineering

25

An odd case

- Map = 100, 000, 100

Run length 0, then 5,

To encode 1000..... (i.e. run length 0)

x = 0 //still takes one bit to represent "0"

y = 0 // run length is 0 long

00 110101

10BitVectorIndex10

Database Engineering

26

Idea (2b) Operations on Compressed bitmaps

Select *
From Faculty
Where Faculty.gender = "F" and
Faculty.FacRank = "asst"

	Bitmap
Asst	110010010001
F	01010...
result	01000...

CounterAsst

0, 0, 2, 2, 3

1, 1,

CounterF

Use counters and arithmetic
in a merge-like operation

10BitVectorIndex10

Database Engineering

27

Given Compression, consider:

- attributes with *many* values

→ lot's of zero's in the maps

→ long runs of zeros, occasional 1's

suppose 100's of zero's in a run

- how fast is the merge compared to logical ANDs?
- how fast does index storage grow with added values?

10BitVectorIndex10

Database Engineering

28

Consider case where: value is a key

- n, records, n = m different values
- each and every map, just one "1".

worst case, run length of n-1,

→ space required for index,

$$n * 2 \log_2 n$$

10BitVectorIndex10

Database Engineering

29

What about real numbers?

- Worst-case
 - Every number in a column is unique
 - > but that's the same as if it were a key

10BitVectorIndex10

Database Engineering

30

Internal Implementation

- Map position assignment
 - rows numbered, stored in a special place/table
(position, row-id) // supported with secondary index
- e.g. (1, row-id₁)
(2, row-id₂)
...


10BitVectorIndex10

Database Engineering

31

Internal Implementation

- Storing/locating bitmaps
 - (key, value pair),
 - key = index key value
 - value = bitmaps



10BitVectorIndex10

Database Engineering

32

Internal Implementation

- Update an indexed value in a row,
 - look up bitmaps for old value and new value
 - flip respective bits (may change vector length)
- Insert new row:
 - assign next row number
 - what happens to existing bitmaps?
- Delete row:
 - set previously set bit to 0
 - “retire” row number (until compaction)

10BitVectorIndex10

Database Engineering

33

How powerful?

MS SQL Server, PostgreSQL
and
IBM DB2

- will build bitmap indexes, dynamically, to execute a query
- then thrown the index out

10BitVectorIndex10

Database Engineering

34

Near Future

- Columns stores

10BitVectorIndex10


Database Engineering

35

Row Storage Column Storage

[(later relates to): horizontal partitioning]

vertical partitioning]



Vertical Partitioning

Part of databases
called “Column-
Stores”

16 Transactions & Recovery

Database Systems

36

Column Stores

- From research
 - MonetDB (open source) [2002]
 - H-store and C-store ~[2005]
- To commercial practice
 - Sybase IQ [1995],
 - bought by SAP...
 - SAP HANA [2008],
 - main-memory, cluster, derived from Sybase IQ
 - HP Vertica
 - [Vertica founded 2005, forked open source C-store fork]
 - Sisense
 - a Tableau competitor, offers similar function but on multiple terabytes on a desktop