


Storage Models

Objectives:

- Detail three storage models,
 - Row storage (conventional)
 - Column Storage
 - Parallel/Distributed Key-Value stores (NoSQL)
- And a first look at parallel/distributed database structure

Row Storage classic, conventional


- That's what we've been doing
 - Row is both a logical and physical unit
- Primary key,
 - Basis of sorting rows on data blocks
 - Search key for the primary index




Horizontal Partition of Row Storage

- First introduction to parallel databases

Shared Nothing [DB] Architecture aka *A cluster*





- Low cost, *commodity*, servers
 - Connected by a network
- Won (at least mind share)

Shared Nothing – Advantages

- Economical: uses commodity hardware
 - Rack mounted servers
- Most scalable
 - Minimizes interference by minimizing resource sharing
 - Memory and I/O bandwidth and capacity grow with the number of compute nodes.


16 Transactions & Recovery

Database Systems

6


Maybe later in the semester

- Shared Memory,



Dan Sudou - CSEPS44 Fall 2010

- Shared Disk



Dan Sudou - CSEPS44 Fall 2010

- My favorite: “Weird Machines”

Other Architectures...

people think they lost

but they actually live,
(they have a low profile)
they could come back

New Idea 1: Partition Data

- Partition Data: to split the storage of a table across the servers.
 - Horizontal Partitioning
 - Vertical Partitioning

16 Transactions & Recovery

Database Systems

9

3 Methods of Horizontal Partitioning

- Round Robin
- Hash
- Range

Horizontal Data Partitioning

- Relation R split into P chunks R_0, \dots, R_{P-1} , stored at the P nodes

Let t_j be a tuple in chunk R_i

Let a_j be the value or set of values
for an attribute or set of attribute for all t_j


16 Transactions & Recovery

Database Systems

11

Horizontal Data Partitioning

- Relation R split into P chunks R_0, \dots, R_{P-1} , stored at the P nodes



16 Transactions & Recovery

Database Systems

12

Round Robin Partitioning

- Relation R split into P chunks R_0, \dots, R_{P-1} , stored at the P nodes
- **Round robin**: tuple t_i to chunk $(i \bmod P)$
 - Like dealing cards from a deck of cards
 - All chunks the same size, $(+/- 1)$
 - If "*the deck*" is randomized, the chunks are randomized

16 Transactions & Recovery

Database Systems

13

Hash Partitioning

- **Hash based partitioning on attribute A**:
 - Tuple t to chunk $h(t.A) \bmod P$
- Hash the value(s) in the tuple to some integer
- That integer maps to a processor number
 - If attribute values are randomized, the chunks are randomized, and roughly the same size

16 Transactions & Recovery

Database Systems

14

Range Partitioning

- **Range based partitioning on attribute A**:
 - Tuple t to chunk i if $v_{i-1} < t.A < v_i$
- Usually the DBA specifies the ranges (v_{i-1}, v_i)


16 Transactions & Recovery

Database Systems

15

Horizontal Data Partitioning

- Relation R split into P chunks R_0, \dots, R_{P-1} , stored at the P nodes
- **Round robin**: tuple t_i to chunk $(i \bmod P)$
- **Hash based partitioning on attribute A**:
 - Tuple t to chunk $h(t.A) \bmod P$
- **Range based partitioning on attribute A**:
 - Tuple t to chunk i if $v_{i-1} < t.A < v_i$



16 Transactions & Recovery

Database Systems

16

Parallel Selection

Compute $\sigma_{A=v}(R)$, or $\sigma_{v1 < A < v2}(R)$

- On a conventional database: cost = $B(R)$
- Q: What is the cost on a parallel database with P processors ?
 - Round robin
 - Hash partitioned
 - Range partitioned

16 Transactions & Recovery

Database Systems

17

Selection – Round Robin Partitioning

Compute $\sigma_{A=v}(R)$, or $\sigma_{v1 < A < v2}(R)$

- Q: What is the cost on a parallel database with P processors ?
- Round robin: all servers do the work
 - Parallel time = $B(R)/P$; total work = $B(R)$
 - Good load balance but need to read all the data

16 Transactions & Recovery

Database Systems

18

Selection - Hash Partitioning

Compute $\sigma_{A=v}(R)$, or $\sigma_{v1 < A < v2}(R)$

- Hash:
 - $\sigma_{A=v}(R)$: Parallel time = total work = $B(R)/P$
- Query is on the hashed attribute

Is A the primary key? Discuss implications

16 Transactions & Recovery

Database Systems

19

Range Partitioning

- Range: one server only
 - Parallel time = total work = $B(R)$ // Discuss
 - Works well for range predicates but suffers from data skew

16 Transactions & Recovery

Database Systems

20

Parallel Selection

- Q: What is the cost on a parallel database with P processors ?
- Round robin: all servers do the work
 - Parallel time = $B(R)/P$; total work = $B(R)$
 - Good load balance but needs to read all the data
- Hash:
 - $\sigma_{A=v}(R)$: Parallel time = total work = $B(R)/P$
 - $\sigma_{A \in [v1, v2]}(R)$: Parallel time = $B(R)/P$; total work = $B(R)$
- Range: one server only
 - Parallel time : total work // Discuss
 - Works well for range predicates but suffers from data skew

16 Transactions & Recovery

Database Systems

21

The Column Store Storage Model

Column Stores: The SQL entry to NoSQL Big Data

1. Column stores **are not** NoSQL, but you will see that in marketing
 - They *do* support Big Data
 - Marketing people don't get modas ponens


NoSQL → Big data

Column Store → Big data

So, marketing says -

Column Store → NoSQL **Not**, and they are being called out on it.

Vertical Partitioning → Column Store



- *Row index (think in terms of an array) is not necessarily the primary key.
- No assurance a declared primary key embody the essential sequential property need to create a bitmap.

The Good

- Storage Space...

- Use Bitmap index (only) as data storage
- Use **Compressed** Bitmap index (only) as data storage

Last Name	First Name	E-mail
1	1	1
2	2	2
3	3	3

The Bad: Consider Insert(new_row)

Last Name	First Name	E-mail	Street
1	1	1	1
2	2	2	2
3	3	3	3

How many disk accesses (writes)?

- Horizontal: [as few as] 1
- Vertical: at least the number of columns
... and then there is the transaction log

So:

- Transactional workload or analytic workload?

Star Schema: (a.k.a. OLAP schema)
SSBM benchmark derived from TPC-H
Current: “dimensional data modeling”

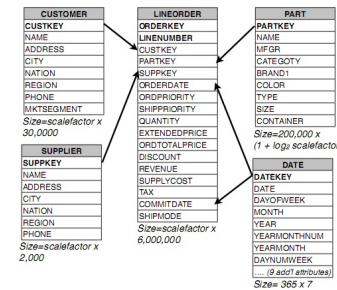


Figure 1: Schema of the SSBM Benchmark

Figure taken from [1]

The Very Good

```
SELECT Sum(S.sales)
FROM sales_table S, time T, location L
WHERE L.city = 'Austin' AND
      T.year = 2019 AND
      S.locationkey = L.pk_loc AND
      S.timekey = T.pk_time
```

- How many columns must be read to process query? 7
 - L.city, L.pk_loc
 - T.year, T.pk_time
 - S.locationkey, S.timekey, S.sales
- In contemporary systems these are stored as compressed bitmaps.
- Compare to amount of data to read in a conventional horizontal store (?)

What actually get's read?

```
SELECT Sum(S.sales)
FROM sales_table S, time T, location L
WHERE L.city = 'Austin' AND
      T.year = 2019 AND
      S.locationkey = L.pk_loc AND
      S.timekey = T.pk_time
```

- L.city, just the bitmap for 'Austin'
- T.year, just the bitmap for 2019

But to do the explicit joins we'll need all values (bitmaps) for

- S.locationkey, L.pk_loc
- S.timekey, T.pk_time

Similarly to compute the output Sum(S.sales), we will need the bitmaps for all values

Improved I/O not without cost:

```
SELECT Sum(S.sales)
FROM sales_table S, time T, location L
WHERE L.city = 'Austin' AND
      T.year = 2019 AND
      S.locationkey = L.pk_loc AND
      S.timekey = T.pk_time
```

What needs to be joined?

- Seemingly unavoidable
 - S.locationkey = L.pk_loc
 - S.timekey = T.pk_time

But to actually compute the result

- For S, the sequential index of S.locationkey, S.timekey, S.sales,
 - must all be the same.
 - Each may take on many values
 - (many) x (many) x (many)
 - In bitmap can be done linear time AND
 - If Not in bitmap, but sorted, linear time merge
- Similarly,
 - T.year, T.pk_time
 - L.city, L.pk_loc
 - But just 1, value, Austin and 2019 have to match index with the other argument (from the same table)
 - → AND the bitmaps
- Hence,
 - 2 obvious joins in the SQL query
 - + 4 operations to assemble the output

Consider

```
SELECT * ** all columns of the three tables
FROM sales_table S, time T, location L
WHERE L.city = 'Austin' AND
      T.year = 2019 AND
      S.locationkey = L.pk_loc AND
      S.timekey = T.pk_time
```

- Reading all columns, all values
- Reassemble columns into rows

Thus, Column Store vs. Row Store Tradeoff

For a given query

- Column store may require many fewer disk block reads than a horizontal store.
- Column store requires computation not needed by a horizontal store to assemble output.
- Note: More column reads → more work to assemble output.

Column Stores

- From research
 - MonetDB (open source) [2002]
 - H-store and C-store ~[2005]
- To commercial practice
 - Sybase IQ [1995], bought by SAP... evolved to:
 - SAP HANA [2008], main-memory, cluster
 - HP Vertica [Vertica founded 2005 from C-store fork]
 - Sisense // a Tableau competitor, offers similar function but on multiple terabytes on a desktop
- “Proof” column store \equiv SQL RDBMS
 - MariaDB (MySQL fork) *Column Store version*
 - SQL Server (starting, 2016 V13), *columnstore indexes*

16 Transactions & Recovery


Database System SQL Server 2016 (13.x), columnstore indexes

34

NoSQL Key Value Stores

1. Parallel/Distributed Storage
2. Notion of numbered Logical and Physical processors
(which will be ignored until near the last slide)
3. Each processor is replicated many times
 - Replicates contain the same data

NoSQL



1: Introduction


Data Management & Engineering

36

Chord Protocol


[<http://www.inf.ed.ac.uk/teaching/courses/ip/chord-desc.html>]

The Chord protocol: given a key, determine the node responsible for storing the key's value



Chord assigns hash keys to nodes in a way that doesn't need to change much as nodes join and leave the system. [SIGCOMM '01 paper by Stoica et al.](#)

NoSQL



1: Introduction

Data Management & Engineering

38

Key, Value model

Setname: {(key1, value1), (key2, value2) (key3, value3)...

- Key_i Typically any string, but also OID (object id)
- store(key_i, value_i) will store (key_i, value_i) replicates in processor Chord(key_i)
 - the key is explicitly stored as its often also data
 - Data type of value_i depends on the system, but can be anything
 - Json document
 - A nested set of (key, value) pairs


1: Introduction

Data Management & Engineering

39

Hashing Shards, Replicate for Durability

- Hashing



1: Introduction

Data Management & Engineering

40

Distributed File System

- Hadoop Distributed File System (HDFS)
 - pages/blocks ~~16~~ or 64 or 128 Mbytes
 - pages are compressed for storage
 - pages are replicated for fault tolerance
 - quorum consistency is the basis of transactions

13 Join Operators

Database Management & Engineering

41

NOSQL: Not as Different as “they” Would Like You to Believe

- E.g.
 - name{ (k1, dan), (k2, bob), (k3, bruce)}
 - salary{ (k1, \$10), (k2, \$1), (k3, \$1000)}
 - title { (k1, professor), (k2, fry cook), (k3, chairman)}
- Employee

Id	Name	Salary	Title
1	dan	10	professor
2	bob	1	fry cook
3	bruce	1000	chairman
...			

1: Introduction

Data Management & Engineering

42

Look familiar?

- E.g.
 - name{ (k1, dan), (k2, bob), (k3, bruce)}
 - salary{ (k1, \$10), (k2, \$1), (k3, \$1000)}
 - title { (k1, professor), (k2, fry cook), (k3, chairman)}
- Employee

Id	Name	Salary	Title
1	dan	10	professor
2	bob	1	fry cook
3	bruce	1000	chairman
...			

1: Introduction

Data Management & Engineering

43

How about now?

- E.g.
 - Employee.name{ (k1, dan), (k2, bob), (k3, bruce)}
 - Employee.salary{ (k1, \$10), (k2, \$1), (k3, \$1000)}
 - Employee.title { (k1, professor), (k2, fry cook), (k3, chairman)}
- Employee

Id	Name	Salary	Title
1	dan	10	professor
2	bob	1	fry cook
3	bruce	1000	chairman
...			

Id	Name	Salary
1	dan	10
2	bob	1
3	bruce	1000
...		

Id	Title
1	professor
2	fry cook
3	chairman
...	

1: Introduction

Data Management & Engineering

44

But we hashed on k_i

Both vertically partitioned

AND

Horizontally hash partitioned

- E.g.
 - Employee.name{ (k1, dan), (k2, bob), (k3, bruce)}
 - Employee.salary{ (k1, \$10), (k2, \$1), (k3, \$1000)}
 - Employee.title { (k1, professor), (k2, fry cook), (k3, chairman)}

- Employee

Id	Name	Salary	Title
1	dan	10	professor
2	bob	1	fry cook
3	bruce	1000	chairman
...			

Id	Name	Id	Salary
1	dan	1	10
2	bob	2	1
3	bruce	3	1000
...		...	

Id	Title
1	professor
2	fry cook
3	chairman
...	

1: Introduction

Data Management & Engineering

45

What if value is a set of (key, value) pairs?

- e.g.


```
Employee{ (k1, ((name dan), (salary 10), (title professor))),
          (k2, ((name bob), (salary 1), (title fry cook))),
          (k3, ((name bruce), (salary 1000), (title chairman)))
        }
```

- Employee

Id	Name	Salary	Title
1	dan	10	professor
2	bob	1	fry cook
3	bruce	1000	chairman
...			

1: Introduction

Data Management & Engineering

46

What if value is a set of (key, value) pairs?

Hash-based horizontal partitioning

- e.g.


```
Employee{ (k1, ((name dan), (salary 10), (title professor))),
          (k2, ((name bob), (salary 1), (title fry cook))),
          (k3, ((name bruce), (salary 1000), (title chairman)))
        }
```

- Employee

Id	Name	Salary	Title
1	dan	10	professor
2	bob	1	fry cook
3	bruce	1000	chairman
...			

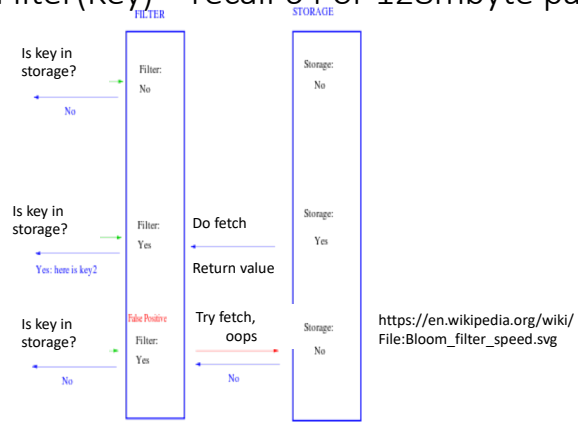
1: Introduction

Data Management & Engineering

47

What about indexing?

Bloom_Filter(Key) – recall 64 or 128mbyte pages



What about secondary indexes?

- Add more Bloom Filters
 - Declare the field used as a key

** pseudo code

CREATE Bloomfilter INDEX foo ON Employee.name(value)

- E.g.
 - Employee.name{ (k1, dan), (k2, bob), (k3, bruce)}
 - Employee.salary{ (k1, \$10), (k2, \$1), (k3, \$1000)}
 - Employee.title { (k1, professor), (k2, fry cook), (k3, chairman)}
- Member(foo, "dan")
 - Executes concurrently on all processors.

Id	Name
1	dan
2	bob
3	bruce
...	

Id	Salary
1	10
2	1
3	1000
...	

Id	Title
1	professor
2	fry cook
3	chairman
...	

1: Introduction

Data Management & Engineering

51

** pseudo code

CREATE Bloomfilter INDEX foo ON
Employee(WHERE subkey = "name")

- e.g.


```
Employee{ (k1, ((name dan), (salary 10), (title professor))),
          (k2, ((name bob), (salary 1), (title fry cook))),
          (k3, ((name bruce), (salary 1000), (title chairman)))
        }
```

- Employee

Id	Name	Salary	Title
1	dan	10	professor
2	bob	1	fry cook
3	bruce	1000	chairman
...			

1: Introduction

Data Management & Engineering

52

Storage Model Summary

Model	Data Storage	Primary Index	Secondary Index
Conventional RDBMS	Rows	B+ tree	B+ tree, plus vendor specific additions
Column Store (RDBMS)	Compressed Bit Maps	N//A	?
Key Value (NoSQL)	Hash partitioned, replicated, large compressed pages	Bloom Filter	Bloom Filter, maybe others