

Introduction to Query Processing

Objectives:

- Learn basic structure/steps in a query execution environment
- An aside: there are *a lot* of things called join

some slides thanks to text authors

11Intro.QueryProcessing

Database Management & Engineering

1

Problem Statement

- Given
 - a SQL query
 - relationally stored data
- 1. What has to happen to return a correct answer?
- 2. What has to happen that answers are returned as fast as possible?

11Intro.QueryProcessing

Database Management & Engineering

2

Character of the Solution

- Executing statements in a computer program.

```
select *
from employees
where employees.salary = 0
```

--> The solution is organized as a compiler/interpreter

11Intro.QueryProcessing

Database Management & Engineering

3

What are the steps?

Compiler:

1. Lex & Parse the input
 1. parse-trees
 2. symbol tables
2. Compile parse-trees to abstract code
3. [optimize abstract code]
4. Compile abstract code to physical code.
5. [optimize physical code]

SQL execution environment

1. Lex & Parse the input
2. Create a logical plan
 - abstract code
3. Optimize, logical-plan
4. Consider physical trade-offs (optimize), generate
→ physical plan (sequence of executable operators, I.e. code)

Difference: Creation of “physical code” is a critical optimization process

11Intro.QueryProcessing

Database Management & Engineering

4

R	A	B	C	S	C	D	E
	a	1	10		10	x	2
	b	1	20		20	y	2
	c	2	10		30	z	2
	d	2	35		40	x	1
	e	3	45		50	y	3

Example

Select B,D

From R,S

Where $R.A = "c" \wedge S.E = 2 \wedge R.C = S.C$

Where $R.A = "c" \wedge S.E = 2 \wedge R.C = S.C$

R	A	B	C	S	C	D	E
	a	1	10		10	x	2
	b	1	20		20	y	2
	c	2	10		30	z	2
	d	2	35		40	x	1
	e	3	45		50	y	3

Answer

B	D
2	x

Some words about *join*

1. Logically, there are many kinds of joins.
2. Each logical join may be implemented by many different algorithms. (*physical operators*)
3. Formal definition of [inner]* join

$$R \bowtie_p S = \sigma_p (R \times S)$$


* Originally, and for a long time there were only *inner* joins

Today

“join” \equiv Any operator combining information from two tables.

- Inner Joins
 - natural join
 - equi join
 - theta join
- Outer Joins
- Anti-Joins
- Semi-joins

Inner Joins

- natural join, , with no argument
 - Equality test on all pairs of columns with the same name.
- Theta join, $R \bowtie_p S$
 - Any relational predicate p (or theta)
- Equi join,
 - A join where the predicate contains only = tests
 - So, all natural joins are equi joins
 - all equi join joins are theta joins

11Intro.QueryProcessing

Database Management & Engineering

10

SQL: All these mean the same thing

```
SELECT *
FROM Persons, Zips
INNER JOIN Persons ON Zips
Persons.zipcode = Zips.zipcode
```

```
SELECT *
FROM Persons, Zips
Where Persons.zipcode = Zips.zipcode
```

```
SELECT *
FROM Persons
WHERE (zipcode, city, state) IN
      (SELECT *
       FROM Zips )
```

```
SELECT *
FROM Persons JOIN Zips
ON Persons.zipcode = Zips.zipcode
```

Example Goal

```
Select B,D
```

```
From R,S
```

```
Where
```

```
R.A = "c" ^
```

```
S.E = 2 ^
```

```
R.C=S.C
```

```
output(
```

```
  project((B,D),
```

```
    nestedLoopJoin("C=C",
```

```
      select("A = "c"", tableScan(R))
```

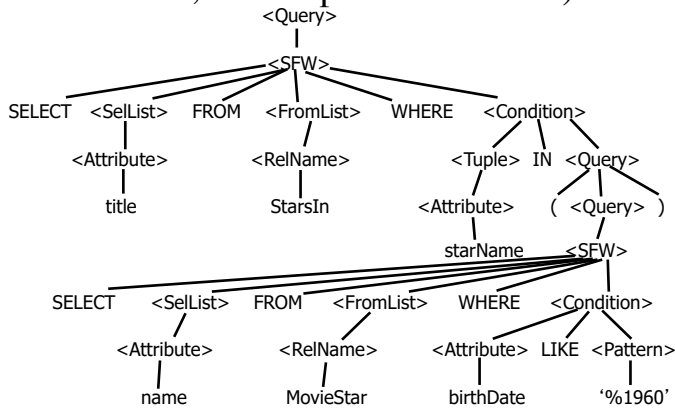
```
      select("E = 2", tableScan(S))))
```

11Intro.QueryProcessing

Database Management & Engineering

12

First step: Parse (standard compiler course stuff; not responsible for this)




11Intro.QueryProcessing

Database Management & Engineering

13

Logical Plan for the Above:




11Intro.QueryProcessing

Database Management & Engineering

14

Probably the Real Starting Point




OR: $\Pi_{B,D} [\sigma_{R.A='c' \wedge S.E=2 \wedge R.C=S.C} (RXS)]$

11Intro.QueryProcessing

Database Management & Engineering

15

We say "we pushed the selects down"



11Intro.QueryProcessing

Database Management & Engineering


16

Generate a logical plan

People version:

$\Pi_{B,D} [\sigma_{R.A="c" \wedge S.E=2 \wedge R.C=S.C} (R \bowtie S)]$

Machine version:



11Intro.QueryProcessing

Database Management & Engineering

17


Logical Optimization

Given

- axioms and identities of relational algebra
- E.g.
 - $(R \bowtie S) = (S \bowtie R)$
 - pushing selects
 - If all columns in p are in R

$$\sigma_p (R \bowtie S) \rightarrow [\sigma_p (R)] \bowtie S$$

Manipulate logical plan



11Intro.QueryProcessing

Database Management & Engineering

18

Big, but not too big, library of transformations, choice is important.

- $R \bowtie S \rightarrow S \bowtie R$ (logically equiv. but not physically)
- $\sigma_{p1 \wedge p2} (R) \rightarrow \sigma_{p1} [\sigma_{p2} (R)]$
- $\sigma_p (R \bowtie S) \rightarrow [\sigma_p (R)] \bowtie S$
- $\pi_x [\sigma_p (R)] \rightarrow \pi_x \{ \sigma_p [\pi_{xz} (R)] \}$

11Intro.QueryProcessing

Database Management & Engineering

19

Physical Plan Generation

Replace logical operators with physical operators, while considering the cost of alternatives

- Hard to separate optimization from generation process.

11Intro.QueryProcessing

Database Management & Engineering

20

Suppose there is an index on R.A?

```
Select B,D
From R,S
Where R.A = "c" ^ S.E = 2 ^ R.C=S.C
```

If it's the primary index, no thinking

```
output(
  project((B,D),
    nestedLoopJoin("C=C",
select("A = 'c'", tableScan(R))
    primaryIndexScan(c,R)
    select("E = 2", tableScan(S))))
```

If it's a secondary index (remember city = "Austin")
~~select("A = 'c'", tableScan(R))~~ or secondaryIndexScan(c,R) ?

Better compute the cost of the alternatives.

11Intro.QueryProcessing

Database Management & Engineering

21

Other components of cost

- size of the relation
 - base relation : a table stored in the database
 - relation: any table, stored or computed
- Given an operator, a predicate a relation
 - (e.g. R.city = "austin")

How big is the computed relation?

11Intro.QueryProcessing

Database Management & Engineering

22

In this next section we will cover

- details of implementing individual database operators.
- how to estimate the cost of the operators
 - cost(operator, predicates, input parameters)
 - returns
 - cost values, I/O, CPU
 - input parameters of the relation for the next cost
- search methods for assigning the final choices.

11Intro.QueryProcessing

Database Management & Engineering

23