

## Merge-Join

First we will consider equijoins on primary index keys.

```
create table customers (cid ... primary key
                        ... )
create table orders (oid ...,
                    cid ... primary key,
                    ...)
```

13 Join Operators

Database Management &amp; Engineering

1

## Merge-Join

equijoins on primary index keys.

```
create table customers (cid ... primary key
                        ... )
create table orders (oid ...,
                    cid ... primary key,
                    ...)
```

- Data is sorted & clustered on disk... by primary key
  - So how does it come off of the disk?
    - HDDs
    - SSDs
- I call this “**The Golden Case**” for merge-join

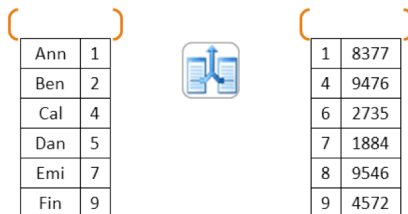
13 Join Operators

Database Management &amp; Engineering

2

## Merge Join

Thanks:  
<https://sqlity.net/en/1480/a-join-a-day-the-sort-merge-join/>



Ann	1
Ben	2
Cal	4
Dan	5
Emi	7
Fin	9

1	8377
4	9476
6	2735
7	1884
8	9546
9	4572

## Careful:

- Above works for join arguments being unique.  
 //no duplicate values in the join argument
- Modify algorithm to allow for duplicate values  
 // left as an exercise

13 Join Operators

Database Management &amp; Engineering

4

## Merge-Join Pseudo Code

*Assume block structure is buried in iterator abstraction*

```
Merge-Join(relation R, relation S, attr rjk, attr sjk ){
  r = open(R, 1) // read a block of R
  s = open(S, 1) // read a block of S
  until( r = EOF or s = EOF) {
    if (r.rjk = s.rjk) {output(r,s), r = next(r)}
    if (r.rjk < s.rjk) {r = next(r)}
    if (r.rjk > s.rjk) {s = next(s)}
  }
}
```

## Cost of Merge-Join

$B(R) + B(S)$  // per text

What happens in real life?

- Hint: consider the affine cost model

What if an equijoin, but not on index key?

Is there a secondary index on the join argument?

Do I want to use the secondary index?


## Assume No Index

Recall two phase sorting


- for all practical purposes relations can be sorted in, 2 reads + 1 write
  - (final write doesn't count, by convention)
- cost to sort  $R = 3 B(R)$  // why 3 and not 4?

## Optimizer will Transform the Plan

### • Logical Plan



### • Physical Plan



13 Join Operators

Database Management &amp; Engineering

9

## Merge-Sort Cost – No Index

$$5(B(R) + B(S))$$

- 3x to sort
- 1x (did have to write the sorted results out)
- 1x for merge join

May feel high – but it is very predictable

13 Join Operators

Database Management &amp; Engineering

10

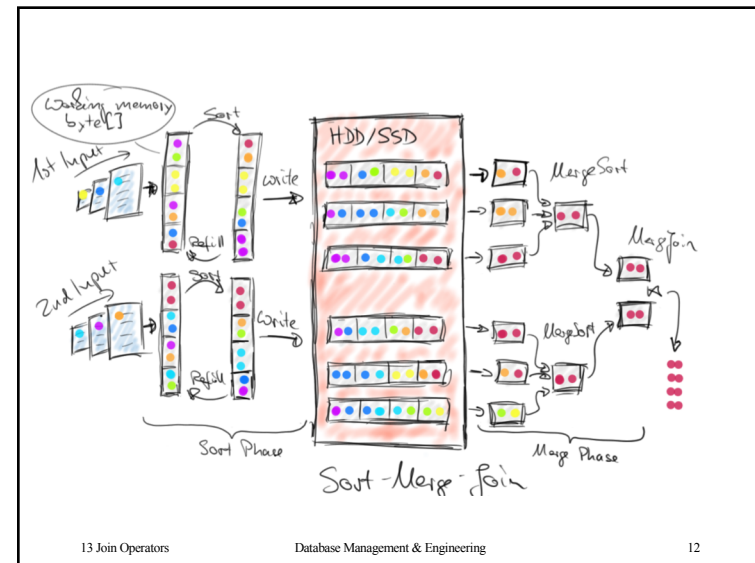
## Many Opportunities to be Clever

- what if 1 relation fits in memory?
  - sort is  $1 \cdot B(R)$
- integrate merge-join with second phase of two phase sort,
  - $3(B(R) + B(S))$

13 Join Operators

Database Management &amp; Engineering

11



13 Join Operators

Database Management &amp; Engineering

12