

## Introduction Parallel Databases

Objectives:

- 1) Sources of Parallelism
- 2) Hardware/Software Architectures
- 3) Effectiveness

Slide thanks: many sources, including Suciu  
Database Systems

1

## DBs Provide Many Opportunities for Parallelism

- I/O Parallelism // RAID or simply many disks
- Interquery Parallelism // Transactions
- Intraquery Parallelism // operations in a plan
  - pipelined,
  - simply independent, Interoperation Parallelism
- Intraoperation Parallelism

[Silberschatz et. al.]

16 Transactions & Recovery

Database Systems

2


## Hardware Configurations

16 Transactions & Recovery

Database Systems

3

## Shared Nothing (DB) Clusters




## Shared Nothing – Advantages

- Most scalable
  - Minimizes interference by minimizing resource sharing
  - Memory and I/O bandwidth and capacity grow with the number of compute nodes.
- Can use commodity hardware

## Other Hardware Organizations


- Shared memory
- Shared disk
- Logic per disk [head] and other weird computers.

## Shared Memory



Dan Suciu – CSEP544 Fall 2010

## Shared Disk



Dan Suciu – CSEP544 Fall 2010

## Weird Machines

- Logic, i.e. a CPU per disk [head]
  - push select and project into disk controller
  - // reduces I/O bandwidth
- “Database machines: an idea whose time has passed? A critique of the future of database machines”
  - Haran Boral and David Dewitt, 1989
  - Core point: must use commodity hardware
- “Query Processing on Smart SSDs”
  - Park et.al. and Dewitt
  - <http://sites.computer.org/debull/A14june/p19.pdf>

## Parallel DBMSs

- **Goal**
  - Improve performance by executing multiple operations in parallel
- **Key benefit**
  - Cheaper to scale than relying on a single increasingly more powerful processor
- **Key challenge**
  - Ensure overhead and contention do not kill performance

## Performance Metrics for Parallel DBMSs

- **Speedup**
  - More processors → higher speed
  - Individual queries should run faster
  - Should do more transactions per second (TPS)
- **Scaleup** // sequential execution time
  - More processors → can process more data
  - **Batch scaleup**
    - Same query on larger input data should take the same time
  - **Transaction scaleup**
    - N-times as many TPS on N-times larger database
    - But each transaction typically remains small


## Speed-up and Amdahl's law

- Speed up, S

$$S = \frac{T_{seq}}{T_{par}}$$

// sequential execution time  
// parallel execution time

## Linear v.s. Non-linear Speedup



16 Transactions & Recovery

Database Systems

13

## Challenges to Linear Speedup and Scaleup

- **Startup cost**
  - Cost of starting an operation on many processors
- **Interference**
  - Contention for resources between processors
- **Skew**
  - Slowest processor becomes the bottleneck

16 Transactions & Recovery

Database Systems

14

## Speed-up and Amdahl's law

- Speed up,  $S$

$$S = \frac{T_{seq}}{T_{par}}$$

- Suppose,  
 $T_{par} = (1 - \alpha) T_{seq} + \alpha T_{seq}/P$  //  $\alpha$  proportion of sequential execution  
 // time that can be parallelized
- Then  
 $S = 1 / ((1 - \alpha) + \alpha/P)$
- Suppose  $\alpha = 0.9$ , number of processors  $P \rightarrow \infty$ , then  $S \rightarrow 10$

16 Transactions & Recovery

Database Systems

15

## How does this relate to HW5a?

16 Transactions & Recovery

Database Systems

16


## How does this relate to HW5a?

- Unless processing starts with data distributed/enabled for parallel access witnessing parallel speed-up is hard.

## How is this so ignored/unknown?


## Current Business as Usual

Cloud storage charges (\$):



## Result: Common Practice

Store Here,



move here only when needed

