

Consistency Constraints (enforcing correct data)

Objectives:

- The DBMS as an active manager of the correctness contents and behavior, enabling more object like behavior.
- Specification of such in SQL

["beer slides" borrowed from Garcia-Molina via Tim Finin]

2a:Definitions

CS386D Database Systems

1

Data Model

“A *data model* is a notation for describing data or information... generally consisting of three parts”

- Structure of the data...
- Operations on the data ...
- Constraints on the data ...

1: Introduction

Data Management & Engineering

2

Context: In the beginning

- Date & Codd: formal mathematical underpinnings of relational databases
 - Many years to make this work:
 - e.g. query optimization - $R \text{ join } S = S \text{ join } R$
 - transactions - row level locks
 - Centered on syntax, correctness.
 - Why?
 - make it work at all.

2a:Definitions

CS386D Database Systems

3

Then what happened?

- We had to make it work in the real world

2a:Definitions

CS386D Database Systems

4

First Step, re: Real World Introducing Primary Index

If primary key is one attribute

```
CREATE TABLE t_test
(a INTEGER PRIMARY KEY,
b VARCHAR(10)
)
```

Values in the primary key
column(s) are unique
→ easy & fast to locate a
particular row

If primary key is a number of attributes it is called a
composite key.

```
CREATE TABLE track(
album CHAR(10),
dsk INTEGER,
posn INTEGER,
song VARCHAR(255),
PRIMARY KEY (album, dsk, posn)
)
```

simple key: one attribute

compound key: multiple attributes

Constraint Syntax and Terminology

If primary key is one attribute

```
CREATE TABLE t_test
(a INTEGER PRIMARY KEY,
b VARCHAR(10)
)
```

If primary key is a number of attributes it is called a
composite key.

```
CREATE TABLE track(
album CHAR(10),
dsk INTEGER,
posn INTEGER,
song VARCHAR(255),
PRIMARY KEY (album, dsk, posn)
)
```

PRIMARY KEY a.k.a. a "primary key
constraint"


- if the constraint is on a single
attribute --> column constraint
- if the constraint is on a multiple
attributes --> table constraint

simple key: one attribute

compound key: multiple attributes

Inventor-Invention, Object Model UML Class Diagram Logical Model


A list of inventions, each with their list of inventors




a 1 to many relationship

a *logical* model

Inventor-Invention, Physical Data Model (UML syntax)



Inventor-Invention, Physical Data Model (UML syntax)



Tool dependent:

- PK, FK, PFK labels on the diagram
- Name munging in the creation of key column names
 - E.g. PK_Invention_ID
 - FK_Invention_PK_Invention_ID

2a:Definitions

CS386D Database Systems

9

Inventor - Invention as DB Tables

Invention	
iid	name
1	structure
2	sequencing_machine
3	expression_chips

Inventor		
iid	first	last
1	Francis	Crick
1	James	Watson
1	Rosalyn	Franklin
1	Maurice	Wilkins
2	Lee	Hood
3	David	Botstein

Who invented the sequencing machine?

```

Select last
From Invention, Inventor
Where

```

```

Invention.name="sequencing_machine"
and
Invention.iid = Inventor.iid

```

foreign key implements
1 to many relation

2a:Definitions

CS386D Database Systems

10

Of Foreign Keys, Pointers and Surrogates

Surrogate: An instance of a data type that can be used to locate a disk block or a row in a database. Can be a key or row id (RID)
- informally, but *never* formally, a pointer.

2a:Definitions

CS386D Database Systems

11

The Problem

- Relational representations are semantically challenged.
 - Formal completeness results that show it is possible to compute/represent “rich” concepts.
 - Complicated structures are not *expressed* elementally (usage, like chemistry)
 - As object programming arose, issues became more obvious (painful)

2a:Definitions

CS386D Database Systems

12

Solution, two parts

- DBMS provides facilities that support richer ideas of database consistency.
 - Constraints, Views (declarative)
 - Triggers (procedural constructs)
- Modeling tools integrate object tools and *compile* the representation to relational form.

2a:Definitions

CS386D Database Systems

13

The Constraints

<https://www.tutorialspoint.com/sql/sql-constraints.htm>

- NOT NULL - column cannot have NULL value.
- DEFAULT – Provides a default value for a column when none is specified.
- UNIQUE – all values in a column are different.
- PRIMARY Key – Uniquely identifies each row/record in a database table.
- FOREIGN Key –
- CHECK Constraint – The CHECK constraint ensures that all the values in a column satisfies certain conditions.

2a:Definitions

CS386D Database Systems

14

Foreign Key: References

- In relation R
 - "attribute A references S (B)"
 - values in the A column of R must
 - uniquely appear
 - in the B column of relation S .
 - (Oracle) B must be declared the primary key for S .

2a:Definitions

CS386D Database Systems

15

Why declare this to the database?

- declare == (inform and enforce)
- performance
 - join example - on the board
 - correctness

2a:Definitions

CS386D Database Systems

16

Example

- CREATE TABLE Beers (
 name CHAR(20) PRIMARY KEY,
 manf CHAR(20)
);
CREATE TABLE Sells (
 bar CHAR(20),
 beer CHAR(20) REFERENCES Beers(name),
 price REAL
);

2a:Definitions

CS386D Database Systems

17

Alternative

- Add another element declaring the foreign key, as:
CREATE TABLE Sells (
 bar CHAR(20),
 beer CHAR(20),
 price REAL,
 FOREIGN KEY beer REFERENCES
 Beers(name)
);
- Extra element essential if the foreign key is more than one attribute.

2a:Definitions

CS386D Database Systems

18

Exceptions

- Two ways:
 - 1. Insert a Sells tuple referring to a nonexistent beer.
 - Always rejected.
 - 2. Delete or update a Beers tuple that has a beer value some Sells tuples refer to.
 - a) Default: reject.
 - b) Cascading delete: Ripple changes to referring Sells tuple.

2a:Definitions

CS386D Database Systems

19

Example

- Delete "Bud." Cascade deletes all Sells tuples that mention Bud.
- Update "Bud" → "Budweiser." Change all Sells tuples with "Bud" in beer column to be "Budweiser."

2a:Definitions

CS386D Database Systems

20

What Happens On a Foreign Key Exception

- c) Set Null : Change referring tuples to have NULL in referring components.
- Example
- Delete "Bud." Set-null makes all Sells tuples with "Bud" in the beer component have NULL there.
- Update "Bud" → "Budweiser." Same change.

2a:Definitions

CS386D Database Systems

21

Selecting a Policy

- Add ON [DELETE, UPDATE] [CASCADE, SET NULL] to declaration of foreign key.

2a:Definitions

CS386D Database Systems

22

Selecting a Policy (II)

- Example


```
CREATE TABLE Sells (
  bar CHAR(20),
  beer CHAR(20),
  price REAL,
  FOREIGN KEY beer REFERENCES
  Beers(name)
  ON DELETE SET NULL
  ON UPDATE CASCADE
);
```

2a:Definitions

CS386D Database Systems

23

Selecting a Policy (III)

- "Correct" policy is a design decision.
 - E.g., what does it mean if a beer goes away?
 - What if a beer changes its name?

2a:Definitions

CS386D Database Systems

24

Attribute-Based Checks

- Follow an attribute by a condition that must hold for that attribute in each tuple of its relation.
- Form: CHECK (condition).
 - Condition may involve the checked attribute.
 - Other attributes and relations may be involved, but only in subqueries.
 - Oracle 7.3.2: No subqueries allowed in condition.

2a:Definitions

CS386D Database Systems

25

Attribute-Based Checks

- Condition is checked only when the associated attribute changes (i.e., an insert or update occurs).

2a:Definitions

CS386D Database Systems

26

Example

- CREATE TABLE Sells (
 - bar CHAR(20),
 - beer CHAR(20) CHECK(
 - beer IN (SELECT name
 - FROM Beers) // much like fk constraint
 -),
 - price REAL CHECK(
 - price <= 5.00
 -)
-);

2a:Definitions

CS386D Database Systems

27

Attribute-Based Checks (III)

- Check on beer is like a foreign-key constraint, except
 - The check occurs only when we add a tuple or change the beer in an existing tuple, not when we delete a tuple from Beers.

2a:Definitions

CS386D Database Systems

28

Tuple/Table-Based Checks

- Separate element of table declaration.
- Form: like attribute-based check.
- But condition can refer to any attribute of the relation.
 - Or to other relations/attributes in subqueries.
 - Again: Oracle 7.3.2 forbids the use of subqueries.

2a:Definitions

CS386D Database Systems

29

Example

- Only Joe's Bar can sell beer for more than \$5.

```
CREATE TABLE Sells (
  bar CHAR(20),
  beer CHAR(20),
  price REAL,
  CHECK(bar = 'Joe's Bar'
        OR price <= 5.00)
);
```

2a:Definitions

CS386D Database Systems

30

Triggers

- Often called event-condition-action rules.
- *Event* = a class of changes in the DB, e.g., "insert into Beers."
- *Condition* = a test as in a where-clause for whether or not the trigger applies.
- *Action* = one or more SQL statements.

2a:Definitions

CS386D Database Systems

31

Example

- Whenever we insert a new tuple into Sells, make sure the beer mentioned is also mentioned in Beers, and insert it (with a null manufacturer) if not.

2a:Definitions

CS386D Database Systems

32

Example (II)

- CREATE TRIGGER BeerTrig
AFTER INSERT ON Sells
FOR EACH ROW
WHEN(new.beer NOT IN (SELECT name FROM Beers))
BEGIN
INSERT INTO Beers(name)
VALUES(:new.beer);
END;


2a:Definitions

CS386D Database Systems

33

How can the DBMS support this? (internal architecture)

- All updates *funnel* through the transaction manager




2a:Definitions

CS386D Database Systems

34

Funnel is implemented in the SQL Engine

- SQL interpreter executes
 - insert, delete, update
 - *commit*
- Commit is a clear boundary
 - actions
 - before
 - after



2a:Definitions

CS386D Database Systems

35

Triggers

- Standardized in SQL3
- Prior to SQL3, vendor specific syntax
(vestigial compatibility exists)
- Action part, of ECA is a procedural program
 - every vendor has their own procedural extensions. e.g. PSQL, T-SQL,...

2a:Definitions

CS386D Database Systems

36

Example (II)

- Sells(bar, beer, price)
 CREATE OR REPLACE TRIGGER BeerTrig
 AFTER INSERT ON Sells
 FOR EACH ROW
 WHEN(new.beer NOT IN (SELECT name FROM Beers))
 BEGIN
 INSERT INTO Beers(name)
 VALUES(:new.beer);
 END;
 .
 Run

2a:Definitions

CS386D Database Systems

37

Options

- Can omit OR REPLACE. Effect is that it is an error if a trigger of this name exists.
- AFTER can be BEFORE.
- INSERT can be DELETE or UPDATE OF < attribute > ON.
- FOR EACH ROW can be omitted, with an important effect: the action is done once for the relation(s) consisting of all changes.

2a:Definitions

CS386D Database Systems

38

Notes

- There are two special variables new and old, representing the new and old tuple in the change.
 - old makes no sense in an insert, and new makes no sense in a delete.
- Oracle:
 - in WHEN we use new and old without a colon, but in actions, a preceding colon is needed.

2a:Definitions

CS386D Database Systems

39

Oracle Notes (II)

- The action is a PL/SQL statement.
 - Simplest form: surround one or more SQL statements with BEGIN and END.
 - However, select-from-where has a limited form.
- Dot and run cause the definition of the trigger to be stored in the database. (stored procedure)
 - Oracle triggers are elements of the database, like tables or views.

2a:Definitions

CS386D Database Systems

40

Example

- Maintain a list of all the bars that raise their price
for some beer by more than \$1.
Sells(bar, beer, price)

2a:Definitions

CS386D Database Systems

41

Example (II)

- CREATE TRIGGER PriceTrig
AFTER UPDATE OF price ON Sells
FOR EACH ROW
WHEN(new.price > old.price + 1.00)
BEGIN
INSERT INTO RipoffBars
VALUES(:new.bar);
END;
.
run

2a:Definitions

CS386D Database Systems

42