

Datalog (2) Evaluating Datalog Programs

10 Datalog I

Database Systems

1

Consider

```
CREATE VIEW A AS
  SELECT B.b1 AS a1, B.b2 AS a2
  FROM B
  ** no where clause on purpose
```

10 Datalog I

Database Systems

2

Consider

```
CREATE VIEW A AS
  SELECT B.b1 AS a1, B.b2 AS a2
  FROM B
```

As a Datalog rule:

$A(a1, a2) \leftarrow B(a1, a2)$ ** remember **positions**, not names

10 Datalog I

Database Systems

3

Consider

```
CREATE VIEW A AS
  SELECT B.b1 AS a1, B.b2 AS a2 FROM B

CREATE VIEW B AS
  SELECT C.c1 AS b1, C.c2 AS b2 FROM C

CREATE VIEW C AS
  SELECT D.d1 AS c1, D.d2 FROM D

CREATE TABLE D (d1 integer, d2 integer);
```

10 Datalog I

Database Systems

4

If D contains (8, 16)

```
CREATE VIEW A AS
  SELECT B.b1 AS a1, B.b2 AS a2 FROM B
```

```
CREATE VIEW B AS
  SELECT C.c1 AS b1, C.c2 AS b2 FROM C
```

```
CREATE VIEW C AS
  SELECT D.d1 AS c1, D.d2 FROM D
```

```
CREATE TABLE D (d1 integer, d2 integer);
```

What does `SELECT * FROM A`
Return?

10 Datalog I

Database Systems

5

Make a plan, compute: $A = \{(8, 16)\}$

What does `SELECT * FROM A`
Return?

```
CREATE VIEW A AS
  SELECT B.b1 AS a1, B.b2 AS a2 FROM B
```

```
CREATE VIEW B AS
  SELECT C.c1 AS b1, C.c2 AS b2 FROM C
```

```
CREATE VIEW C AS
  SELECT D.d1 AS c1, D.d2 FROM D
```


```
CREATE TABLE D (d1 integer, d2 integer);
```

A top-down evaluation

10 Datalog I

Database Systems

6



As a Datalog program

$$A(a1, a2) \leftarrow B(a1, a2)$$

$$B(b1, b2) \leftarrow C(b1, b2)$$

$$C(c1, c2) \leftarrow D(c1, c2)$$

$$D = \{(8, 16)\}$$

10 Datalog I

Database Systems

7



Then, Suppose:

INSERT INTO D VALUES(32, 64)

As a Datalog program

$$A(a1, a2) \leftarrow B(a1, a2)$$

$$B(b1, b2) \leftarrow C(b1, b2)$$

$$C(c1, c2) \leftarrow D(c1, c2)$$


$$D = \{(8, 16), (32, 64)\}$$

We can re-execute the plan

10 Datalog I

Database Systems

8



But what if

```
CREATE TABLE D (d1 integer, d2 integer);
INSERT INTO D VALUES(8,16);
```

```
CREATE MATERIALIZED VIEW C AS
  SELECT D.d1 AS c1, D.d2 FROM D
```

```
CREATE MATERIALIZED VIEW B AS
  SELECT C.c1 AS b1, C.c2 AS b2 FROM C
```

```
CREATE MATERIALIZED VIEW A AS
  SELECT B.b1 AS a1, B.b2 AS a2 FROM B
```

10 Datalog I

Database Systems

9

And Then INSERT INTO D VALUES(32, 64)

$$D = \{(8, 16)\} \cup (32, 64)$$

Incrementally Maintain Materialized View

10 Datalog I

Database Systems

10

From past lecture, but incrementally
maintained view had a join

$$(R \cup S) \mid X \mid T = (R \mid X \mid T) \cup (S \mid X \mid T)$$

So, if View, $V = (S \mid X \mid T)$

◆ Insert Into S, rows R, i.e. $S' = (R \cup S)$

◆ New view contents, V' , $V' = S' \mid X \mid T$

$$\begin{aligned} V' &= (R \cup S) \mid X \mid T \\ &= (R \mid X \mid T) \cup (S \mid X \mid T) \\ &= (R \mid X \mid T) \cup V \end{aligned}$$

7:Semantics/Triggers/V
iews

Database Systems

11

So:

If $D = \{(8, 16)\}$

$$C(c1, c2) \leftarrow D(c1, c2)$$

$$B(b1, b2) \leftarrow C(b1, b2)$$

$$A(a1, a2) \leftarrow B(a1, a2)$$

$$C = \{(8, 16)\}$$

$$B = \{(8, 16)\}$$

$$A = \{(8, 16)\}$$

$$D = \{(8, 16), (32, 64)\}$$

A
|
Π
|
B
|
Π
|
C
|
Π
|
D

10 Datalog I

Database Systems

12

Further:

And If $D = \{(8, 16)\} \cup \{(32, 64)\}$ A
 $C(c1, c2) \leftarrow D(c1, c2)$ $C = \{(8, 16)\} \cup \{(32, 64)\}$ Π
 $B(b1, b2) \leftarrow C(b1, b2)$ $B = \{(8, 16)\} \cup \{(32, 64)\}$ $|$
 $A(a1, a2) \leftarrow B(a1, a2)$ $A = \{(8, 16)\} \cup \{(32, 64)\}$ B

By using incremental view maintenance method:

(32, 64) was added to C, but (8, 16) not recomputed

Then

(32, 64) was added to B, but (8, 16) not recomputed

...

A *bottom-up* evaluation

10 Datalog I

Database Systems

A
 Π
 $|$
B
 Π
 $|$
C
 Π
 $|$
D

13

In Datalog, this is called semi-naïve evaluation

And If $D = \{(8, 16)\} \cup \{(32, 64)\}$ A
 $C(c1, c2) \leftarrow D(c1, c2)$ $C = \{(8, 16)\} \cup \{(32, 64)\}$ Π
 $B(b1, b2) \leftarrow C(b1, b2)$ $B = \{(8, 16)\} \cup \{(32, 64)\}$ $|$
 $A(a1, a2) \leftarrow B(a1, a2)$ $A = \{(8, 16)\} \cup \{(32, 64)\}$ B

Next steps, engage with:

- Only think about Datalog,
and not to trivial programs
- Deal with recursion
- Deal with recursion and negation
(*need stratification*)

10 Datalog I

Database Systems

A
 Π
 $|$
B
 Π
 $|$
C
 Π
 $|$
D

14

Example: Datalog Program

- ◆ Using EDB **Sells(bar, beer, price)** and **Beers(name, manf)**, find the manufacturers of beers Joe doesn't sell.

JoeSells(b) <- Sells(' Joe' 's Bar', b, p)

Answer(m) <- Beers(b,m)

AND NOT JoeSells(b)

10 Datalog I

Database Systems

15

Dependency Graph

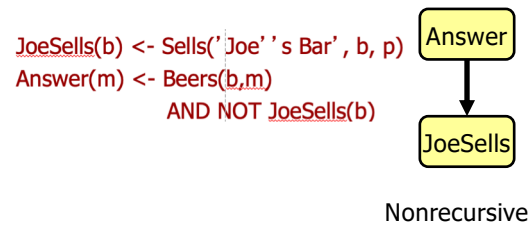
- ◆ Form a *dependency graph* whose nodes = IDB predicates.
- ◆ Arc $X \rightarrow Y$ if and only if there is a rule with X in the head and Y in the body.

10 Datalog I

Database Systems

16

Example: Dependency Graphs



10 Datalog I

Database Systems

17

Evaluating Non-recursive Datalog Programs

- ◆ Build the IDB dependency graph
- ◆ Topologically sort the graph
- ◆ Evaluate the IDB predicates,
 - ▶ Topological sort →
 - all the predicates in the body of its rules have already been evaluated.
 - If an IDB predicate has more than one rule, each rule contributes tuples to its relation.

10 Datalog I

Database Systems

18

Definition of Recursion

- ◆ Form a *dependency graph* whose nodes = IDB predicates.
- ◆ Arc $X \rightarrow Y$ if and only if there is a rule with X in the head and Y in the body.
- ◆ Cycle = recursion; no cycle = no recursion.

10 Datalog I

Database Systems

19

Expressive Power of Datalog

- ◆ Without recursion, Datalog can express all and only the queries of core relational algebra.
 - ▶ The same as SQL select-from-where, without aggregation and grouping.
- ◆ But with recursion, Datalog can express more than these languages.
- ◆ Yet still not Turing-complete.

10 Datalog I

Database Systems

20

Recursive Example


- ◆ EDB: $\text{Par}(c,p) = p$ is a parent of c .
- ◆ Generalized cousins: people with common ancestors one or more generations back:
 $\text{Sib}(x,y) \leftarrow \text{Par}(x,p) \text{ AND } \text{Par}(y,p) \text{ AND } x \neq y$
 $\text{Cousin}(x,y) \leftarrow \text{Sib}(x,y)$
 $\text{Cousin}(x,y) \leftarrow \text{Par}(x,xp) \text{ AND } \text{Par}(y,yp) \text{ AND } \text{Cousin}(xp,yp)$

10 Datalog I

Database Systems

21

Example: Dependency Graphs



Recursive

$\text{Sib}(x,y) \leftarrow \text{Par}(x,p) \text{ AND } \text{Par}(y,p) \text{ AND } x \neq y$
 $\text{Cousin}(x,y) \leftarrow \text{Sib}(x,y)$
 $\text{Cousin}(x,y) \leftarrow \text{Par}(x,xp) \text{ AND } \text{Par}(y,yp) \text{ AND } \text{Cousin}(xp,yp)$

10 Datalog I

Database Systems

22

Consider Subset of Cousin Program

```

(1) sibling(X,Y) :- parent(X,Z) & parent(Y,Z) & X≠Y.
(2) cousin(X,Y) :- parent(X,Xp) & parent(Y,Yp) &
    sibling(Xp,Yp).

```

10 Datalog I

Database Systems


23

Consider Subset of Cousin Program

```

(1) sibling(X,Y) :- parent(X,Z) & parent(Y,Z) & X≠Y.
(2) cousin(X,Y) :- parent(X,Xp) & parent(Y,Yp) &
    sibling(Xp,Yp).

```




10 Datalog I

Database Systems

24

The “Naïve” Evaluation Algorithm



Evaluating Recursive Rules

- ◆ The following works when there is no negation:
 1. Start by assuming all IDB relations are empty.
 2. Repeatedly evaluate the rules using the EDB and the previous IDB, to get a new IDB.
 3. End when no change to IDB. (fixed-point)

10 Datalog I

Database Systems

26

Example: Evaluation of Cousin

- ◆ We'll proceed in rounds to infer Sib facts (red) and Cousin facts (green).
- ◆ Remember the rules:

$\text{Sib}(x,y) \leftarrow \text{Par}(x,p) \text{ AND } \text{Par}(y,p) \text{ AND } x \neq y$
 $\text{Cousin}(x,y) \leftarrow \text{Sib}(x,y)$
 $\text{Cousin}(x,y) \leftarrow \text{Par}(x,xp) \text{ AND } \text{Par}(y,yp) \text{ AND } \text{Cousin}(xp,yp)$

10 Datalog I


Database Systems

27

Par Data: Parent Above Child

(1) $\text{sibling}(X,Y) :- \text{parent}(X,Z) \ \& \ \text{parent}(Y,Z) \ \& \ X \neq Y.$
 (2) $\text{cousin}(X,Y) :- \text{parent}(X,Xp) \ \& \ \text{parent}(Y,Yp) \ \& \ \text{sibling}(Xp,Yp).$
 (3) $\text{cousin}(X,Y) :- \text{parent}(X,Xp) \ \& \ \text{parent}(Y,Yp) \ \& \ \text{cousin}(Xp,Yp).$

Red: sibling rows
Green: cousin rows



Seminaive Evaluation

- ◆ Since the EDB never changes, on each round we only get new IDB tuples ... if we use at least one IDB tuple that was obtained on the previous round.
- ◆ Saves work; lets us avoid rediscovering *most* known facts.
 - ▶ A fact could still be derived in a second way, (a different chain of rules)

10 Datalog I

Database Systems

29

Recursion Plus Negation

- ◆ “Naïve” evaluation doesn’t work when there are negated subgoals.
- ◆ In fact, negation wrapped in a recursion makes no sense in general.
- ◆ Even when recursion and negation are separate, we can have ambiguity about the correct IDB relations.

10 Datalog I

Database Systems

30

Problematic Recursive Negation

$P(x) \leftarrow Q(x) \text{ AND NOT } P(x)$
 EDB: $Q(1), Q(2)$

Initial: $P = \{ \}$

Round 1: $P = \{(1), (2)\}$

Round 2: $P = \{ \}$

Round 3: $P = \{(1), (2)\}, \text{ etc., etc. ...}$

10 Datalog I

Database Systems

31

Stratified Negation

- ◆ Stratification is a constraint usually placed on Datalog with recursion and negation.
- ◆ It rules out negation wrapped inside recursion.
- ◆ Gives the sensible IDB relations when negation and recursion are separate.

10 Datalog I

Database Systems

32

Strata

- ◆ Intuitively, the *stratum* of an IDB predicate P is the maximum number of negations that can be applied to an IDB predicate used in evaluating P .
- ◆ Stratified negation = “finite strata.”
- ◆ Notice in $P(x) \leftarrow Q(x) \text{ AND NOT } P(x)$, we can negate P an infinite number of times deriving $P(x)$.

10 Datalog I

Database Systems

33

Stratum Graph

- ◆ To formalize strata use the *stratum graph* :
 - ▶ Nodes = IDB predicates.
 - ▶ Arc $A \rightarrow B$ if predicate A depends on B .
 - ▶ Label this arc “–” if the B subgoal is negated.

10 Datalog I

Database Systems

34

Stratified Negation Definition

- ◆ The *stratum* of a node (predicate) is the maximum number of – arcs on a path leading from that node.
- ◆ A Datalog program is *stratified* if all its IDB predicates have finite strata.


10 Datalog I

Database Systems

35

The Problematic Example

$P(x) \leftarrow Q(x) \text{ AND NOT } P(x)$



Recursion cycles through a negation

10 Datalog I

Database Systems

36

Another Example

- ◆ EDB = Source(x), Target(x), Arc(x,y).
- ◆ Rules for “targets not reached from any source”:
 - Reach(x) <- Source(x)
 - Reach(x) <- Reach(y) AND Arc(y,x)
 - NoReach(x) <- Target(x)
 - AND NOT Reach(x)

10 Datalog I

Database Systems

37

** Reachability

- ◆ Reachability, as a fundamental capability of graph databases
- A core capability used to promote them

10 Datalog I

Database Systems

38

** Reachability

- ◆ Reachability, as a fundamental capability of graph databases
- A core capability used to promote them
- ◆ Covid-19 contact mapping:
 - ▶ Graph data:
 - vertex is a person,
 - Edge connects two people if they were in physical proximity and, the edge is labeled with the property (key) of the time and date of the contact (value). E.g. (contact, "April 5, 12:07pm")
 - ▶ Who were possible sources of John's Covid-19?
 - All vertices reachable from John in the last 7 days.
 - ▶ Who is the most likely source?
 - What is the shortest path to symptomatic person?

10 Datalog I

Database Systems


- ◆ Each person is represented by vertex.
- ◆ An edge exists from p1 to p2, if

10 Datalog I

Database Systems

40

The Stratum Graph



10 Datalog I

Database Systems

41

Models

- ◆ A *model* is a choice of IDB relations that, with the given EDB relations makes all rules true regardless of what values are substituted for the variables.
 - ▶ Remember: a rule is true whenever its body is false.
 - ▶ But if the body is true, then the head must be true as well.

10 Datalog I

Database Systems

42

Minimal Models

- ◆ When there is no negation, a Datalog program has a unique minimal model (one that does not contain any other model).
- ◆ But with negation, there can be several minimal models.
- ◆ The stratified model is the one that “makes sense.”

10 Datalog I

Database Systems

43

The Stratified Model

- ◆ When the Datalog program is stratified, we can evaluate IDB predicates lowest-stratum-first.
- ◆ Once evaluated, treat it as EDB for higher strata.


10 Datalog I

Database Systems

44

Example: Multiple Models --- (1)

$\text{Reach}(x) \leftarrow \text{Source}(x)$
 $\text{Reach}(x) \leftarrow \text{Reach}(y) \text{ AND } \text{Arc}(y,x)$
 $\text{NoReach}(x) \leftarrow \text{Target}(x) \text{ AND NOT } \text{Reach}(x)$



10 Datalog I

Database Systems

45

Assumption

- ◆ When the logic is stratified, the stratified model is the one that “makes sense.”
- ◆ This principle is used in SQL-99 recursion --- the stratified model is defined to be the correct query result.

10 Datalog I

Database Systems

46

Example: Multiple Models --- (2)

$\text{Reach}(x) \leftarrow \text{Source}(x)$
 $\text{Reach}(x) \leftarrow \text{Reach}(y) \text{ AND } \text{Arc}(y,x)$
 $\text{NoReach}(x) \leftarrow \text{Target}(x) \text{ AND NOT } \text{Reach}(x)$



Another model! Reach(1), Reach(2),
Reach(3), Reach(4); NoReach is empty.

10 Datalog I

Database Systems

47