

## Enforcing Serializability (implementing concurrency control)

- Objectives:
- Locking
  - two phase-locking protocol (with simple locking)
  - escalating locks (read/shared vs write/exclusive)
  - update locks
  - multigranular locks (tuple, block, table level locks)

18Locking

Data Management & Engineering

1

## How to enforce serializable schedules?

### *Option 1: Optimistic protocol*

- run system, record  $P(S)$ ;
- at commit, check  $P(S)$  for cycles
  - if no cycles, commit
  - otherwise, abort offending transactions and restart.

Under what circumstances would this make sense?


18Locking

Data Management & Engineering

2

## How to enforce serializable schedules?

### *Option 2: prevent $P(S)$ cycles from occurring*



18Locking

Data Management & Engineering


3

## A locking protocol

Two new actions:

lock (exclusive):  $li(A)$

unlock:  $ui(A)$



18Locking

Data Management & Engineering

4

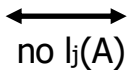
## Rule #1: Well-formed transactions

$T_i$ : ...  $l_i(A)$  ...  $p_i(A)$  ...  $u_i(A)$  ...

Obvious, proper use of lock/unlock.

- Every read or write is preceded by a lock
- Every lock has a corresponding unlock
- Don't ask for a lock twice, not without an intervening unlock:  $l_i(A) l_i(A)$

## Rule #2 Legal scheduler

$S = \dots l_i(A) \dots u_i(A) \dots$   


Blocks a transaction that asks for a lock on a locked object

## Exercise:

- What schedules are legal?  
What transactions are well-formed?

$S1 = l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$   
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

$S2 = l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$   
 $l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$

$S3 = l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$   
 $l_2(B)r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

## Exercise:

- What schedules are legal?  
What transactions are well-formed?

$S1 = l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$   
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

$S2 = l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$   
 $l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$

$S3 = l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$   
 $l_2(A)r_2(A)w_2(A)u_2(A)l_3(B)r_3(B)u_3(B)$

## Exercise:

- What schedules are legal?

What transactions are well-formed?

S1 =  $l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$

$r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$  *illegal*

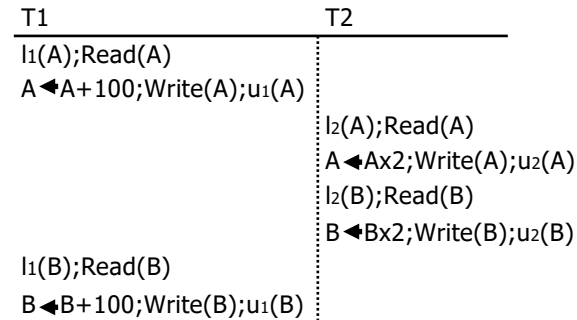
S2 =  $l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$

$l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$  *not well formed*


S3 =  $l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$

$l_2(A)r_2(A)w_2(A)u_2(A)l_3(B)r_3(B)u_3(B)$

## Revisit: Schedule F, add locking



## Schedule F



## Two phase locking (2PL)

for transactions


$T_i = \dots \dots l_i(A) \dots \dots u_i(A) \dots \dots$



- no locks awarded/asked for

after the first unlock

--> collect locks, then release locks




18Locking

Data Management & Engineering

13

## Schedule G




18Locking

Data Management & Engineering

14

## Schedule G

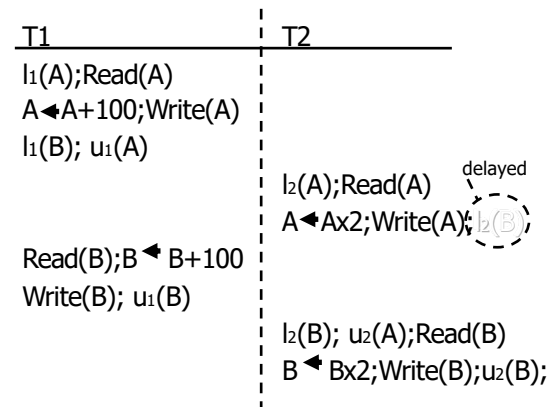


18Locking

Data Management & Engineering

15

## Schedule G




18Locking

Data Management & Engineering

16

## Schedule H (T<sub>2</sub> reversed)



- Assume deadlocked transactions are rolled back
  - They have no effect
  - They do not appear in schedule

E.g., Schedule H =   
 This space intentionally left blank!

Next step:

Show that rules #1,2,3  $\Rightarrow$  conflict-serializable schedules

Conflict rules for l<sub>i</sub>(A), u<sub>i</sub>(A):

- l<sub>i</sub>(A), l<sub>j</sub>(A) conflict
- l<sub>i</sub>(A), u<sub>j</sub>(A) conflict

Note: no conflict < u<sub>i</sub>(A), u<sub>j</sub>(A)>, < l<sub>i</sub>(A), r<sub>j</sub>(A)>, ...

Theorem Rules #1,2,3  $\Rightarrow$  conflict  
(2PL) serializable  
schedule

To help in proof:

Definition  $\text{Shrink}(T_i) = \text{SH}(T_i) =$   
first unlock action of  $T_i$

Lemma

$T_i \rightarrow T_j \text{ in } S \Rightarrow \text{SH}(T_i) <_S \text{SH}(T_j)$

Proof of lemma:

$T_i \rightarrow T_j$  means that

$S = \dots p_i(A) \dots q_j(A) \dots$ ;  $p, q$  conflict

By rules 1,2:

$S = \dots p_i(A) \dots u_i(A) \dots l_j(A) \dots q_j(A) \dots$

By rule 3:  $\text{SH}(T_i) \quad \text{SH}(T_j)$

So,  $\text{SH}(T_i) <_S \text{SH}(T_j)$

Theorem Rules #1,2,3  $\Rightarrow$  conflict  
(2PL) serializable  
schedule

Proof:

(1) Assume  $P(S)$  has cycle

$T_1 \rightarrow T_2 \rightarrow \dots T_n \rightarrow T_1$

(2) By lemma:  $\text{SH}(T_1) < \text{SH}(T_2) < \dots < \text{SH}(T_1)$

(3) Impossible, so  $P(S)$  acyclic


(4)  $\Rightarrow S$  is conflict serializable

- Beyond this simple 2PL protocol, it is all a matter of improving performance and allowing more concurrency....
  - Shared locks
  - Multiple granularity
  - Inserts, deletes and phantoms
  - Other types of C.C. mechanisms

## Shared locks

So far:

$S = \dots l_1(A) \ r_1(A) \ u_1(A) \dots l_2(A) \ r_2(A) \ u_2(A) \dots$



Instead:

$S = \dots ls_1(A) \ r_1(A) \ ls_2(A) \ r_2(A) \dots us_1(A) \ us_2(A)$

## Lock actions

$l-t_i(A)$ : lock A in t mode (t is S or X)

$u-t_i(A)$ : unlock t mode (t is S or X)

Shorthand:

$u_i(A)$ : unlock whatever modes

$T_i$  has locked A

## Rule #1\_ Well formed transactions

$T_i = \dots l-S_1(A) \dots r_1(A) \dots u_1(A) \dots$

$T_i = \dots l-X_1(A) \dots w_1(A) \dots u_1(A) \dots$

Locks on same object are now allowed...

*but - really :*

**escalating** locks on same object allowed

- What about transactions that read and write same object?

Option 1: Request exclusive lock

$T_i = \dots l-X_1(A) \dots r_1(A) \dots w_1(A) \dots u(A) \dots$

- What about transactions that read and write same object?

### Option 2: Upgrade

(E.g., need to read, but don't know if will write...)

$T_i = \dots l-S_i(A) \dots r_1(A) \dots l-X_1(A) \dots w_1(A) \dots u(A) \dots$

Think of  
 - Get 2nd lock on A, or  
 - Drop S, get X lock

18Locking

Data Management & Engineering

29

### Rule #2 Legal scheduler

$S = \dots l-S_i(A) \dots \dots u_i(A) \dots$

↔

no  $l-X_j(A)$

$S = \dots l-X_i(A) \dots \dots u_i(A) \dots$

↔

no  $l-X_j(A)$

no  $l-S_j(A)$

18Locking

Data Management & Engineering

30

### A way to summarize Rule #2

#### Compatibility matrix

		request	
		S	X
held	S	true	false
	X	false	false

18Locking

Data Management & Engineering

31

### Rule # 3 2PL transactions

No change except for upgrades:

- (I) If upgrade gets more locks  
 (e.g.,  $S \rightarrow \{S, X\}$ ) then no change!
- (II) If upgrade releases read (shared)  
 lock (e.g.,  $S \rightarrow X$ )  
 - beginning of shrinking

18Locking

Data Management & Engineering

32



Theorem Rules 1,2,3  $\Rightarrow$  Conf.serializable  
for S/X locks schedules

Proof: similar to X locks case

Detail:

$l-t_i(A), l-r_j(A)$  do not conflict if  $\text{comp}(t,r)$

$l-t_i(A), u-r_j(A)$  do not conflict if  $\text{comp}(t,r)$

Lock types beyond S/X

Examples:


(1) increment lock

(2) update lock

Example (1): increment lock

- Atomic increment action:  $IN_i(A)$   
 $\{\text{Read}(A); A \leftarrow A+k; \text{Write}(A)\}$

- $IN_i(A), IN_j(A)$  do not conflict!



Comp

	S	X	I
S			
X			
I			

Comp

	S	X	I
S	T	F	F
X	F	F	F
I	F	F	T


18Locking

Data Management & Engineering

37

## Update locks

A common deadlock problem with upgrades:



18Locking

Data Management & Engineering

38


## Solution

If  $T_i$  wants to read  $A$  and knows it may later want to write  $A$ , it requests update lock (not shared)

18Locking

Data Management & Engineering

39



18Locking

Data Management & Engineering

40

Comp

Lock already held in

	New request			
	S	X	U	
S	T	F	T	
X	F	F	F	
U	F	F	F	

18Locking

Data Management & Engineering

41

Note: object A may be locked in different modes at the same time...

$S_1 = \dots I-S_1(A) \dots I-S_2(A) \dots I-U_3(A) \dots \left\{ \begin{array}{l} I-S_4(A) \dots ? \\ I- \\ U_4(A) \dots ? \end{array} \right.$

- To grant a lock in mode t, mode t must be compatible with all currently held locks on object

18Locking

Data Management & Engineering

42

How does locking work in practice?

- Every system is different  
(E.g., may not even provide CONFLICT-SERIALIZABLE schedules)
- But here is one (simplified) way ...


18Locking

Data Management & Engineering

43

Sample Locking System:

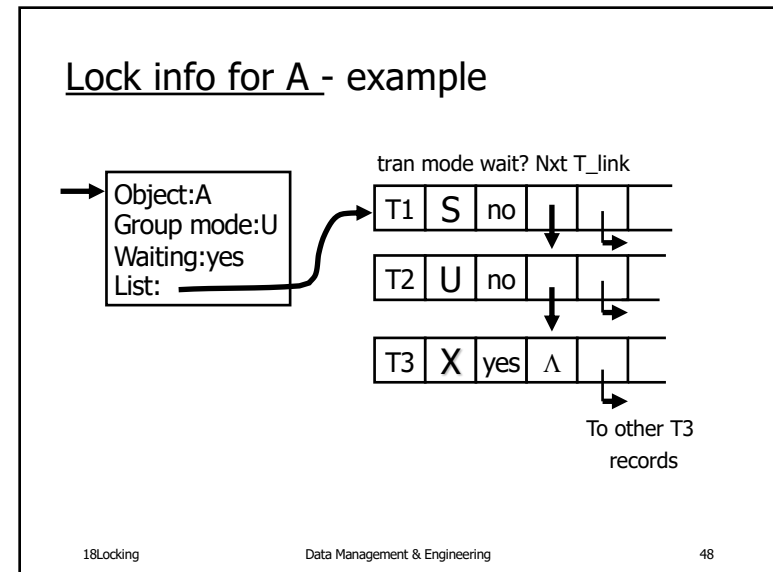
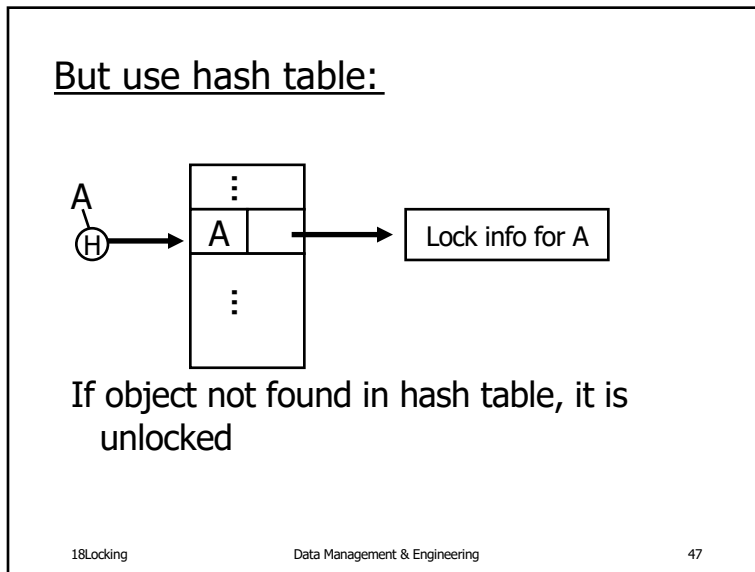
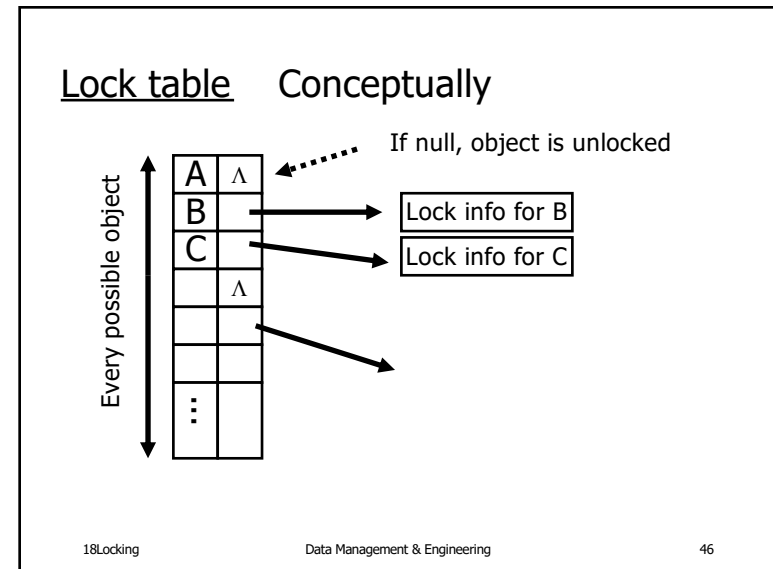
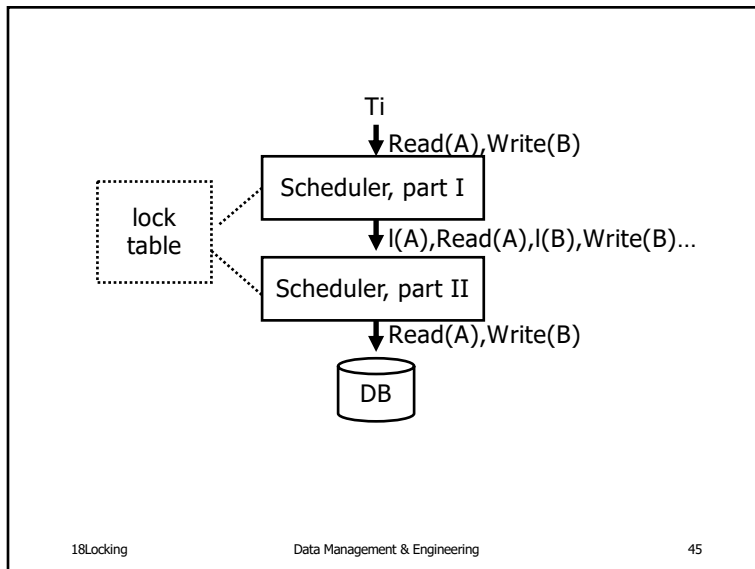
- (1) Don't trust transactions to request/release locks
- (2) Hold all locks until transaction commits



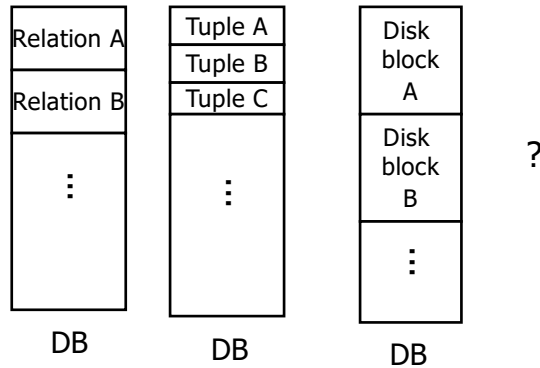
18Locking

Data Management & Engineering

44



What are the objects we lock?



18Locking

Data Management & Engineering

49

- Locking works in any case, but should we choose small or large objects?

- If we lock large objects (e.g., Relations)

- Need few locks
- Low concurrency

- If we lock small objects (e.g., tuples, fields)

- Need more locks
- More concurrency


18Locking

Data Management & Engineering

50

We can have it both ways!!

Ask any janitor to give you the solution...




18Locking

Data Management & Engineering

51

Example




18Locking

Data Management & Engineering

52

## Example



18Locking

Data Management & Engineering

53

## Multiple granularity

Comp

Requestor

		IS	IX	S	SIX	X
Holder	IS					
	IX					
	S					
	SIX					
	X					

18Locking

Data Management & Engineering

54

## Multiple granularity

Comp

Requestor

		IS	IX	S	X
Holder	IS	T	T	T	F
	IX	T	T	F	F
	S	T	F	T	F
	SIX				
	X	F	F	F	F

18Locking


Data Management & Engineering

55

Parent  
locked in

Child can be  
locked in

IS	
IX	
S	
SIX	
X	




18Locking

Data Management & Engineering

56

Parent locked in	Child can be locked in
IS	IS, S
IX	IS, S, IX, X
S	[S, IS] not necessary
X	none



18Locking

Data Management & Engineering

57

## Rules

- (1) Follow multiple granularity comp function
- (2) Lock root of tree first, any mode
- (3) Node Q can be locked by Ti in S or IS only if parent(Q) locked by Ti in IX or IS
- (4) Node Q can be locked by Ti in X,IX only if parent(Q) locked by Ti in IX
- (5) Ti is two-phase
- (6) Ti can unlock node Q only if none of Q's children are locked by Ti


18Locking

Data Management & Engineering

58

## Exercise:

- Can T2 access object f2.2 in X mode?  
What locks will T2 get?




18Locking

Data Management & Engineering

59

## Exercise:

- Can T2 access object f2.2 in X mode?  
What locks will T2 get?




18Locking

Data Management & Engineering

60

### Exercise:

- Can T<sub>2</sub> access object f<sub>3.1</sub> in X mode?  
What locks will T<sub>2</sub> get?




18Locking

Data Management & Engineering

61

### Exercise:

- Can T<sub>2</sub> access object f<sub>2.2</sub> in S mode?  
What locks will T<sub>2</sub> get?



18Locking

Data Management & Engineering

62

## Isolation Levels

### Hard and Soft Transactions in RDBMS

- SQL defines four *isolation levels* = choices about what interactions are allowed by transactions that execute at about the same time.
- Only one level (“serializable”) = ACID transactions.
- Each DBMS implements transactions in its own way.

63

## Choosing the Isolation Level

- Within a transaction, we can say:  
SET TRANSACTION ISOLATION LEVEL *X*  
where *X* =
  - SERIALIZABLE
  - REPEATABLE READ
  - READ COMMITTED
  - READ UNCOMMITTED

64