

Enforcing Serializability Part 2

Objectives:

- Locking – generalizing from simple locks
 - escalating locks (read/shared vs write/exclusive)
 - update locks
 - multigranular locks (tuple, block, table level locks)
- Soft Isolation in SQL
- How to identify deadlock

Left off Last Time:

- Theorem & proof that 2 phase locking allows only conflict serializable schedules.

Theorem Rules #1,2,3 \Rightarrow conflict
(2PL) serializable
schedule

Proof:

(1) Assume P(S) has cycle

$$T_1 \rightarrow T_2 \rightarrow \dots T_n \rightarrow T_1$$

(2) By lemma: $SH(T_1) < SH(T_2) < \dots < SH(T_1)$

(3) Impossible, so P(S) acyclic


(4) \Rightarrow S is conflict serializable

- Beyond this simple 2PL protocol, it is all a matter of improving performance and allowing more concurrency....
 - Shared locks
 - Multiple granularity
 - Inserts, deletes and phantoms
 - Other types of C.C. mechanisms

Shared locks

So far:

$S = \dots l_1(A) \ r_1(A) \ u_1(A) \dots l_2(A) \ r_2(A) \ u_2(A) \dots$



Instead:

$S = \dots ls_1(A) \ r_1(A) \ ls_2(A) \ r_2(A) \dots us_1(A) \ us_2(A)$

Lock actions

$l-t_i(A)$: lock A in t mode (t is S or X)

$u-t_i(A)$: unlock t mode (t is S or X)

Shorthand:

$u_i(A)$: unlock whatever modes

T_i has locked A

Rule #1_ Well formed transactions

$T_i = \dots l-S_1(A) \dots r_1(A) \dots u_1(A) \dots$

$T_i = \dots l-X_1(A) \dots w_1(A) \dots u_1(A) \dots$

Locks on same object are now allowed...

but - really :

escalating locks on same object allowed

- What about transactions that read and write same object?

Option 1: Request exclusive lock

$T_i = \dots l-X_1(A) \dots r_1(A) \dots w_1(A) \dots u(A) \dots$

- What about transactions that read and write same object?

Option 2: Upgrade

(E.g., need to read, but don't know if will write...)

$T_i = \dots l-S_i(A) \dots r_1(A) \dots l-X_1(A) \dots w_1(A) \dots u(A) \dots$

Think of
 - Get 2nd lock on A, or
 - Drop S, get X lock

18Locking

Data Management & Engineering

9

Rule #2 Legal scheduler

$S = \dots l-S_i(A) \dots \dots u_i(A) \dots$

↔

no $l-X_j(A)$

$S = \dots l-X_i(A) \dots \dots u_i(A) \dots$

↔

no $l-X_j(A)$

no $l-S_j(A)$

18Locking

Data Management & Engineering

10

A way to summarize Rule #2

Compatibility matrix

| | | request | |
|------|---|---------|-------|
| | | S | X |
| held | S | true | false |
| | X | false | false |

18Locking

Data Management & Engineering

11

Rule # 3 2PL transactions

No change except for upgrades:

- (I) If upgrade gets more locks
 (e.g., $S \rightarrow \{S, X\}$) then no change!
- (II) If upgrade releases read (shared)
 lock (e.g., $S \rightarrow X$)
 - beginning of shrinking

18Locking

Data Management & Engineering

12

Theorem Rules 1,2,3 \Rightarrow Conf.serializable
for S/X locks schedules

Proof: similar to X locks case

Detail:

$l-t_i(A), l-r_j(A)$ do not conflict if $\text{comp}(t,r)$

$l-t_i(A), u-r_j(A)$ do not conflict if $\text{comp}(t,r)$

Compatibility matrix as a proof tool

| | | request | |
|------|---|---------|-------|
| | | S | X |
| held | S | true | false |
| | X | false | false |

- The compatibility matrix enables,
 - Methodically enumerate, (each cell), the conflicts.
 - Proof skeleton works, just have to argue correctness of the matrix.

Lock types beyond S/X


Examples:

(1) increment lock

(2) update lock

Example (1): increment lock

- Atomic increment action: $IN_i(A)$
 $\{ \text{Read}(A); A \leftarrow A+k; \text{Write}(A) \}$
- $IN_i(A), IN_j(A)$ do not conflict!



Comp

| | S | X | I |
|---|---|---|---|
| S | | | |
| X | | | |
| I | | | |

18Locking

Data Management & Engineering

17

Comp

| | S | X | I |
|---|---|---|---|
| S | T | F | F |
| X | F | F | F |
| I | F | F | T |


18Locking

Data Management & Engineering

18

Update locks

A common deadlock problem with upgrades:



18Locking

Data Management & Engineering

19

Solution

If T_i wants to read A and knows it may later want to write A, it requests update lock (not shared)

18Locking

Data Management & Engineering

20

Comp

New request

Lock already held in

| | S | X | U |
|---|---|---|---|
| S | | | |
| X | | | |
| U | | | |

18Locking

Data Management & Engineering

21

Comp

New request

Lock already held in

| | S | X | U |
|---|---|---|---|
| S | T | F | T |
| X | F | F | F |
| U | F | F | F |

18Locking

Data Management & Engineering

22

Note: object A may be locked in different modes at the same time...

$S_1 = \dots I-S_1(A) \dots I-S_2(A) \dots I-U_3(A) \dots \left\{ \begin{array}{l} I-S_4(A) \dots ? \\ I- \\ U_4(A) \dots ? \end{array} \right.$

- To grant a lock in mode t, mode t must be compatible with all currently held locks on object

18Locking

Data Management & Engineering

23

How does locking work in practice?

- Every system is different
(E.g., may not even provide CONFLICT-SERIALIZABLE schedules)
- But here is one (simplified) way ...

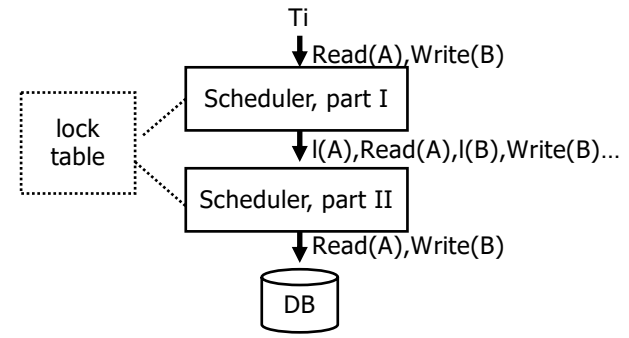
18Locking

Data Management & Engineering


24

Sample Locking System:


- (1) Don't trust transactions to request/release locks
- (2) Hold all locks until transaction commits



Lock table Conceptually




But use hash table:



If object not found in hash table, it is unlocked

Lock info for A - example

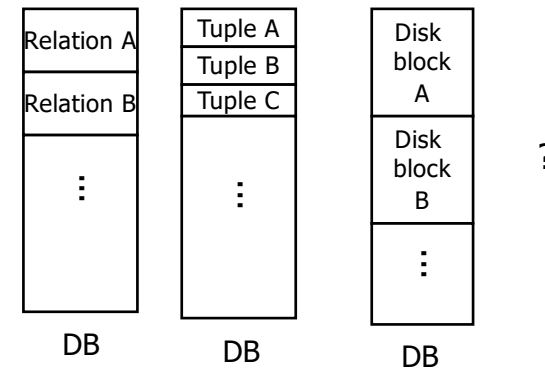


18Locking

Data Management & Engineering

29

What are the objects we lock?



18Locking

Data Management & Engineering

30

- Locking works in any case, but should we choose small or large objects?
- If we lock large objects (e.g., Relations)
 - Need few locks
 - Low concurrency
- If we lock small objects (e.g., tuples, fields)
 - Need more locks
 - More concurrency


18Locking

Data Management & Engineering

31

We can have it both ways!!

Ask any janitor to give you the solution...




18Locking

Data Management & Engineering

32

Example




18Locking

Data Management & Engineering

33

Example



18Locking

Data Management & Engineering

34

Multiple granularity

Comp

Requestor

IS IX S SIX X

| | | | | | | |
|--------|-----|--|--|--|--|--|
| Holder | IS | | | | | |
| | IX | | | | | |
| | S | | | | | |
| | SIX | | | | | |
| | X | | | | | |

18Locking

Data Management & Engineering

35

Multiple granularity

Comp

Requestor

IS IX S X


| | | | | | |
|--------|-----|---|---|---|---|
| Holder | IS | T | T | T | F |
| | IX | T | T | F | F |
| | S | T | F | T | F |
| | SIX | T | F | T | F |
| | X | F | F | F | F |

18Locking

Data Management & Engineering

36

| Parent locked in | Child can be locked in |
|------------------|------------------------|
| IS | |
| IX | |
| S | |
| SIX | |
| X | |




18Locking

Data Management & Engineering

37

| Parent locked in | Child can be locked in |
|------------------|------------------------|
| IS | IS, S |
| IX | IS, S, IX, X |
| S | [S, IS] not necessary |
| X | none |
| | |



18Locking

Data Management & Engineering

38

Rules

- (1) Follow multiple granularity comp function
- (2) Lock root of tree first, any mode
- (3) Node Q can be locked by Ti in S or IS only if parent(Q) locked by Ti in IX or IS
- (4) Node Q can be locked by Ti in X,IX only if parent(Q) locked by Ti in IX
- (5) Ti is two-phase
- (6) Ti can unlock node Q only if none of Q's children are locked by Ti


18Locking

Data Management & Engineering

39

Exercise:

- Can T2 access object f2.2 in X mode?
What locks will T2 get?




18Locking

Data Management & Engineering

40

Exercise:

- Can T2 access object f2.2 in X mode?
What locks will T2 get?




18Locking

Data Management & Engineering

41

Exercise:

- Can T2 access object f3.1 in X mode?
What locks will T2 get?




18Locking

Data Management & Engineering

42

Exercise:

- Can T2 access object f2.2 in S mode?
What locks will T2 get?



18Locking

Data Management & Engineering

43

SQL Isolation Levels

Hard and Soft Transactions in RDBMS

- SQL defines four *isolation levels* = choices about what interactions are allowed by transactions that execute at about the same time.
- Only one level (“serializable”) = ACID transactions.
- Each DBMS implements transactions in its own way.

44

Choosing the Isolation Level

- Within a transaction, we can say:
SET TRANSACTION ISOLATION LEVEL X
where X =
 1. SERIALIZABLE
 2. REPEATABLE READ
 3. READ COMMITTED
 4. READ UNCOMMITTED

45

What About Deadlock?

19LeftOverSouffle

Database Management & Engineering

46

Waits-for-Graph

- Used for dead lock detection.
- Intuitively similar to a precedence graph

But

- edges are populated dynamically.
- Instead of conflicts, edge, (T_i, T_j) exists if T_i is blocked from running by a lock held by T_j
- A cycle in the graph \rightarrow transactions are deadlocked.

19LeftOverSouffle

Database Management & Engineering

47

Waits-for-Graph

- Need a daemon to periodically wake up and check the waits-for-graph for cycles.
- How often?
 - A tradeoff
 - For a large number of transactions, can be expensive.
 - \rightarrow Less often.
 - If there is a deadlock, transactions (customers) are waiting.
 - \rightarrow More often

19LeftOverSouffle

Database Management & Engineering

48