

Final Bits:

- Why Graph Databases are a Growing Force.
- LRU? Not really,
 - In RDBMS its Clock replacement.

Graph Databases and Query Languages...

- Are better Why? How?

Ryan Wolter's Term Project Spring '19

- Implement an Application in Both

SQL and Cypher
And
Compare

The Application – Prologue

- *Very Sadly*, The Application is Proprietary
 - We can't publish.
 - Even in class, I can't reveal much about the application.

The Application

- Application area: Computer Security Forensics
- Practical nature:
 - An actual, labor intensive task
 - Done regularly
 - At the OS command line
 - By a team

19LeftOverSouffle

Database Management & Engineering

5

The Application (cont'd)

- Application area: Computer Security Forensics
- Ryan's insight
 - If the system state(s) examined by the team were represented as data in a database.
 - Entire tasks could be accomplished using a query.
- The data forms a graph!

19LeftOverSouffle

Database Management & Engineering

6

The Experiment

- Developed a use-case document
 - That is: Identified actual tasks to be done.
 - 8 use cases (identifiable tasks).
- Defined schema, SQL and Cypher
- 8 use cases turned into 10 queries
 - 2 use cases took two different queries
- Wrote and ran the queries on Postgres and Neo4j
 - Compared results

19LeftOverSouffle

Database Management & Engineering

7

Qualitative Assessment of the Optimizers

- Neo4j:
 - Has an optimizer that chooses good plans
 - Very primitive,
 - There are explicit instructions on how to write a Cypher query, so optimizer produces better plan.
 - Take the form of,
 - Certain operators, if lexically earlier in the query... will be evaluated earlier in the plan.
 - (developers get to *push down* operations)

19LeftOverSouffle

Database Management & Engineering

8

Qualitative Assessment of the Optimizers

- Postgres:
 - wrt recursive queries
 - Also has, some, sensitivity, to developer syntax

19LeftOverSouffle

Database Management & Engineering

9

Qualitative Assessment of the Optimizers

- Text (string) operations
 - Neo4j does well.
 - SQL uses the LIKE operator
 - Postgres, as is common among RDBMS, does a poor job with LIKE queries.
- With data scraped from non-traditional sources, this has become an important operation.

19LeftOverSouffle

Database Management & Engineering

10

Performance (execution time)

Table 2: Performance Results

Query	Database	Speed
Find exact "all.txt"	Neo4j	1 ms
Find exact "all.txt"	Postgresql	1 ms
Find all *.txt (using :is_root)	Neo4j	4013 ms
Find all *.txt (using WITH)	Neo4j	777 ms
Find all *.txt (using WHERE)	Neo4j	947 ms
Find all *.txt	Postgresql	5949 ms
Find all *.txt ordered limit 10 (executing limit late)	Neo4j	833 ms
Find all *.txt ordered limit 10 (push limit early)	Neo4j	236 ms
Find all *.txt ordered limit 10 (executing limit late)	Postgresql	6008 ms
Find all *.txt ordered limit 10 (push limit early)	Postgresql	32 ms
Find all *.txt ordered limit 1000	Neo4j	636 ms
Find all *.txt ordered limit 1000	Postgresql	56 ms
Find all *.txt ordered limit 10000	Neo4j	636 ms
Find all *.txt ordered limit 10000	Postgresql	5720 ms
Find all *.txt ordered limit 10000 (no recursion)	Postgresql	183 ms

19LeftOverSouffle

Database Management & Engineering

11


Graph Databases? It's not about speed... of the computer

19LeftOverSouffle

Database Management & Engineering

12

Recall Magic Sets Example



Datalog

Rule 1: $\text{Comp}(\text{Part}, \text{Subpt}) \text{ :- } \text{Assembly}(\text{Part}, \text{Subpt}, \text{Qty})$.

Rule 2: $\text{Comp}(\text{Part}, \text{Subpt}) \text{ :- } \text{Assembly}(\text{Part}, \text{Part2}, \text{Qty}), \text{Comp}(\text{Part2}, \text{Subpt})$.

19LeftOverSouffle

Database Management & Engineering

13

Recall Magic Sets Example (2)

Datalog

Rule 1: $\text{Comp}(\text{Part}, \text{Subpt}) \text{ :- } \text{Assembly}(\text{Part}, \text{Subpt}, \text{Qty})$.

Rule 2: $\text{Comp}(\text{Part}, \text{Subpt}) \text{ :- } \text{Assembly}(\text{Part}, \text{Part2}, \text{Qty}), \text{Comp}(\text{Part2}, \text{Subpt})$.

SQL

```
WITH RECURSIVE Comp(Part, Subpt) AS /* Define Comp */
(SELECT A1.Part, A1.Subpt FROM Assembly A1)
UNION
(SELECT A2.Part, C1.Subpt
FROM Assembly A2, Comp C1
WHERE A2.Subpt=C1.Part)
```

$\text{SELECT } * \text{ FROM Comp C2 } /* \text{Returns all parts of a trike } */$

19LeftOverSouffle

Database Management & Engineering

14

Magic Sets Example (3) - Cypher

Cypher

$\text{MATCH} (n:\text{Part})\text{--}[\text{SUBPART_OF}]\text{--}>(p:\text{Part}\{\text{name}:"\text{trike"}\}) \text{RETURN } n.\text{name}$

SQL

```
WITH RECURSIVE Comp(Part, Subpt) AS /* Define Comp */
(SELECT A1.Part, A1.Subpt FROM Assembly A1)
UNION
(SELECT A2.Part, C1.Subpt
FROM Assembly A2, Comp C1
WHERE A2.Subpt=C1.Part)
```

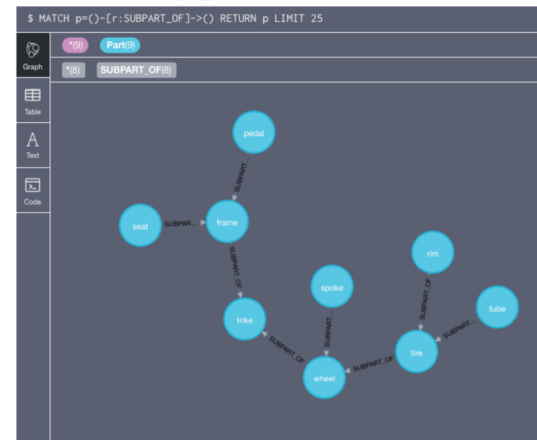
$\text{SELECT } * \text{ FROM Comp C2 } /* \text{Returns all parts of a trike } */$

19LeftOverSouffle

Database Management & Engineering

15

Cypher - IDE



19LeftOverSouffle

16

Real World Example (i.e. Ryan's)

Cypher

```
match(f:file) where f.name = "all.txt"
match(r:is_root)
match p = (r)-[:parent_of*]->(f)
// The reduce notation concatenates the parent nodes to reconstruct the path
return reduce(acc = "/", x IN nodes(p)[1..] | acc + "/" + x.name),
    reduce(acc = 0, x IN nodes(p)[1..] | acc + 1)
```

SQL

```
WITH RECURSIVE filetree AS (
    select file_id, filename, parent_file_id, host, path as path_org,
           filename as path, 0 as depth, parent_file_id as tpid from files
    where filename = 'all.txt'
    UNION
    select ft.file_id, ft.filename, ft.parent_file_id, ft.host,
           ft.path as path_org, f.filename || '/' || ft.path as path,
           ft.depth + 1 as depth, f.parent_file_id as tpid
    from files f
    join filetree ft
    on ft.tpid = f.file_id
```

19LeftOverSouffle

Database Management & Engineering

17

Graph DB: Developers are More Productive

- How important is that?

19LeftOverSouffle

Database Management & Engineering

18

Graph DB: Developers are More Productive

- Assembler
- C
- C++
- Java
- Map Reduce (for parallel programming)

19LeftOverSouffle

Database Management & Engineering

19

Clock Replacement

- A “lower-overhead” approximate implementation of LRU.
- Commonly argued to be good in the face of sequential flooding – (I’ve never bought it)
 - Sequential flooding: What happens, with LRU when repeated scanning R, and, e.g. $B(R) = M(R) + 1$

19LeftOverSouffle

Database Management & Engineering

20

Buffer Replacement Policy slide thanks Joe Hellerstein

- Buffer is chosen for replacement by a *replacement policy*:
 - Least-recently-used (LRU)
 - Most-recently-used (MRU)
 - Also called toss-immediate
 - Clock
- Policy can have big impact on # of I/O's; depends on the *access pattern*.

19LeftOverSouffle

Database Management & Engineering

21

LRU Replacement Policy

- Least Recently Used (LRU)
 - for each page in buffer pool, keep track of time when last *unpinned*
 - replace the frame which has the oldest (earliest) time
 - very common policy: intuitive and simple
 - Works well for repeated accesses to popular pages
- Problems?
- Problem: Sequential flooding
 - LRU + repeated sequential scans.
 - # buffer frames < # pages in file means each page request causes an I/O.
 - Idea: MRU better in this scenario?

19LeftOverSouffle

Database Management & Engineering

22

LRU Replacement Policy

- Least Recently Used (LRU)
 - for each page in buffer pool, keep track of time when last *unpinned*
 - replace the frame which has the oldest (earliest) time
 - very common policy: intuitive and simple
 - Works well for repeated accesses to popular pages
- Problems?
- Problem: Sequential flooding
 - LRU + repeated sequential scans.
 - # buffer frames < # pages in file means each page request causes an I/O.
 - Idea: MRU better in this scenario?


19LeftOverSouffle

Database Management & Engineering

23

“Clock” Replacement Policy

- Associate with the buffers, a cycle of frames
- Per frame, store
 - one *reference bit*
 - *pin count*
- On a buffer access,
 - decrement pin count, until 0
 - when pin count = 0 set reference bit



```

• replacement
do for each page in cycle {
    if (pincount == 0 && ref bit is on)
        turn off ref bit;
    else if (pincount == 0 && ref bit is off)
        choose this page for replacement;
} until a page is chosen;

```

Questions:

How like LRU?

Problems?

19LeftOverSouffle

Database Management & Engineering

Facets of Clock Policy

- Pin buffers by initializing pin count to ∞
- Hack pin count for other purposes,
 - (nested loop join, anticipate how many times the page will be read).
- LRU is *the worst* policy for sequential flooding. Clock by itself isn't that much better, but you can hack pin count.